

Lecture 6:

Convolutional Neural Networks (CNN)

Part 1

Olexandr Isayev

Department of Chemistry, CMU

olexandr@cmu.edu

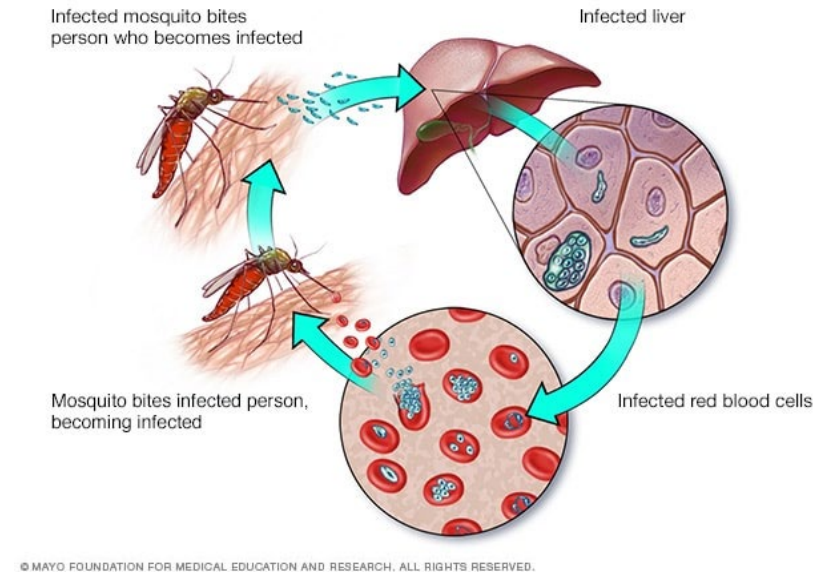
In Class Projects

Please select a problem and dataset

- Upload title and one paragraph summary to canvas.
- I encourage you to use your **domain-specific dataset**
- If not, I am happy to give you one. Just send email or see after class.
- Deadline: **March 3**

HW2: Fighting Malaria with CNNs!

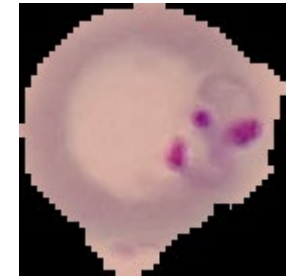
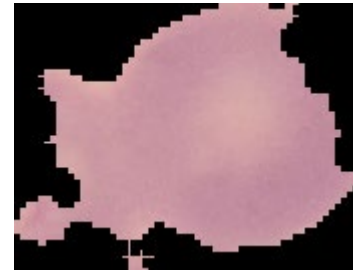
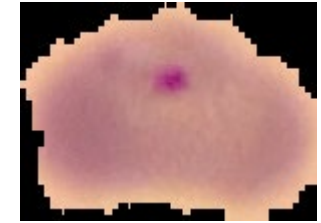
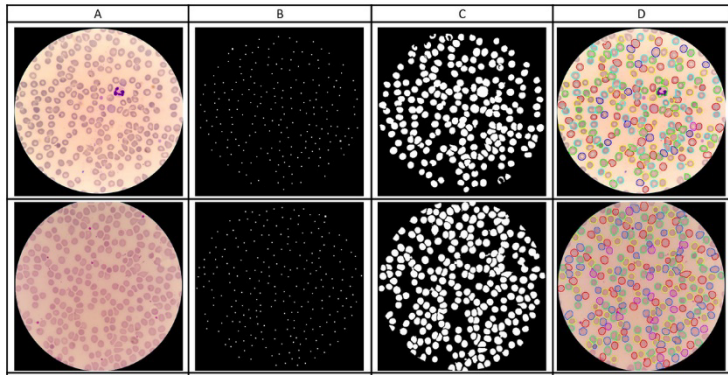
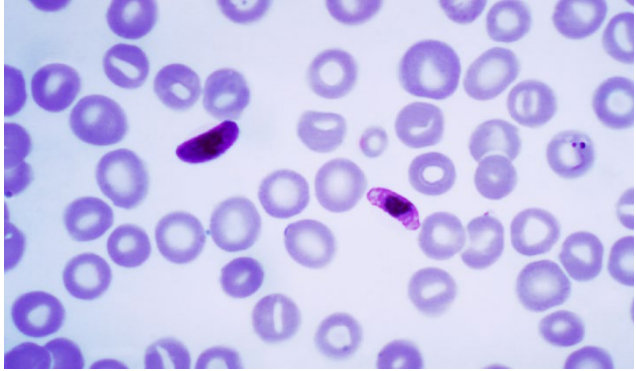
Malaria, a disease caused by protozoan parasites of the genus *Plasmodium*, is not only an acute life threat in many developing countries but also a significant burden on the healthcare system worldwide.



In this homework, you are required to implement a convolution neural network-based machine learning model to predict whether the cell in the given picture is parasitized by the genus *Plasmodium* or not.



Training and Test Data



~25,000 images

Initial blood sample data processing

Your Mission

Kaggle competition

Get data from kaggle.com

Solve the problem

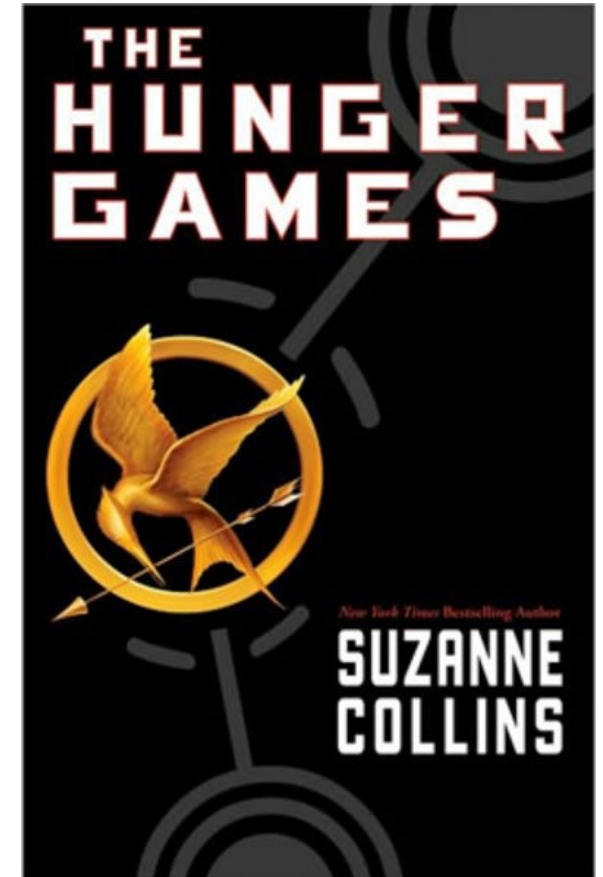
Submit solution for autograding to Kaggle

Submit IPYNB file to canvas

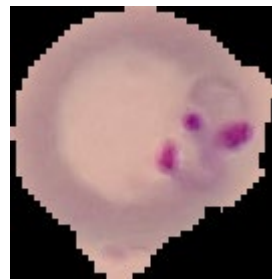
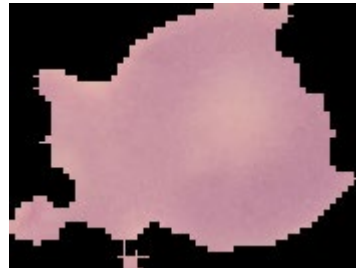
Deadline: **March 11**

Sign-up Link:

<https://www.kaggle.com/t/b7af9038607d45afadb8d22bbdfc0c5d>



Data



x

0
1
0
1
0
0
1
1
1
0
0
1
1
.
.
.
.

y

What you need to do:

- Given the training data: **train_data.zip** and **train_labels.csv**
 - You are allowed to build any type of CNN classification model using PyTorch!
 - You are allowed to use any type of data processing (data augmentation)
 - Please explore training from scratch or finetuning/pertaining existing model .
 - Find the best performing model
 - Use your model to score **test_data.zip**
-
- In this work you will predict the binarized disease state for classification between sick (class 1) and healthy (class 0) blood samples.
 - Performance will be measured using accuracy score (sklearn.metrics.accuracy_score)
 - Overfitting is prevented by using public and private leaderboards (50%-50%).

Kaggle demo

Standard RGB image

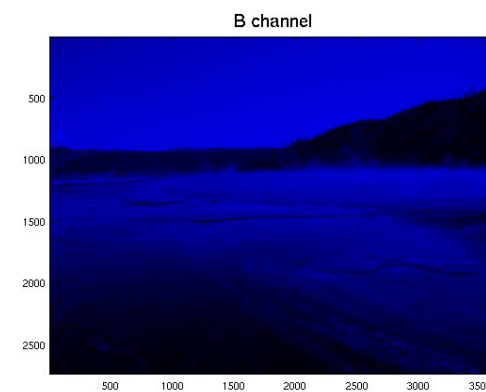
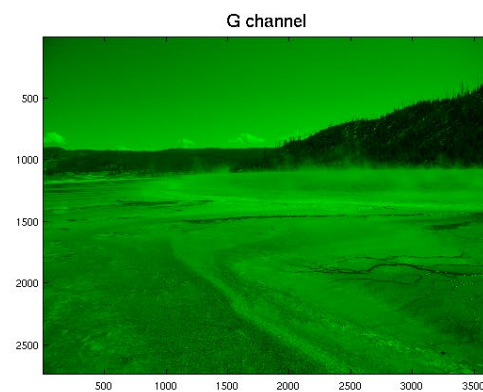
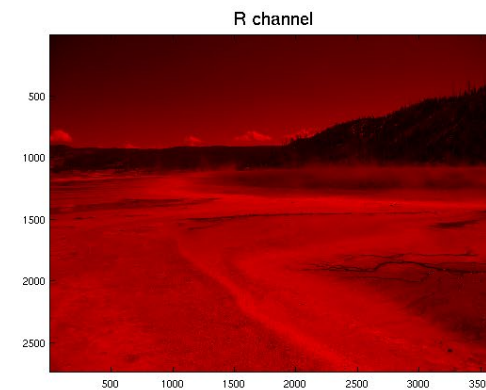
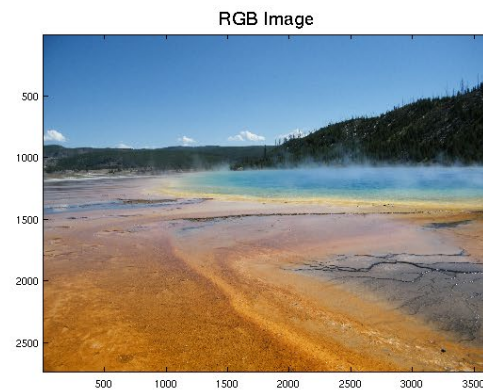
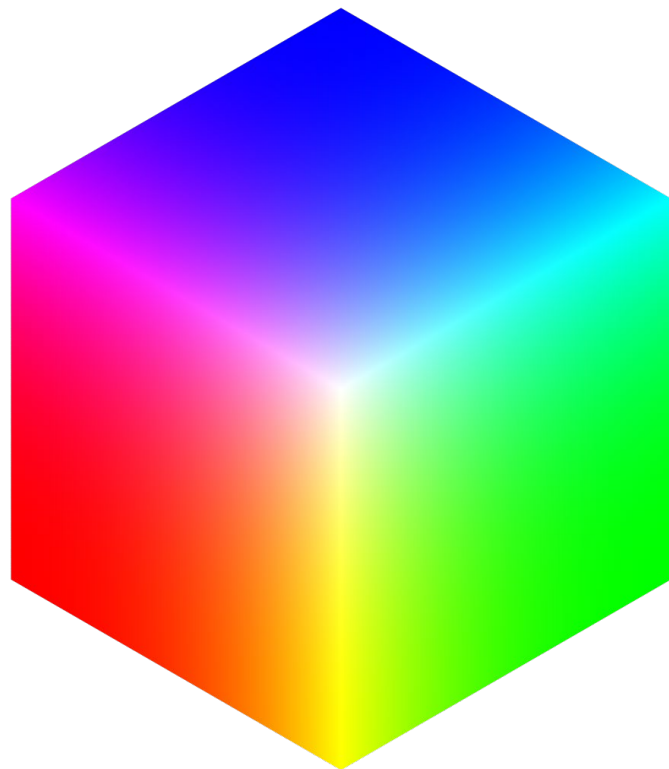


Image features

- We've been basically talking about **detecting features in images**, in a very naïve way.
- Researchers built multiple computer vision techniques to deal with these issues: SIFT, FAST, SURF, BRIEF, etc.
- However, similar problems arose: the detectors where either too general or too over-engineered. Humans were designing these feature detectors, and that made them either too simple or hard to generalize.

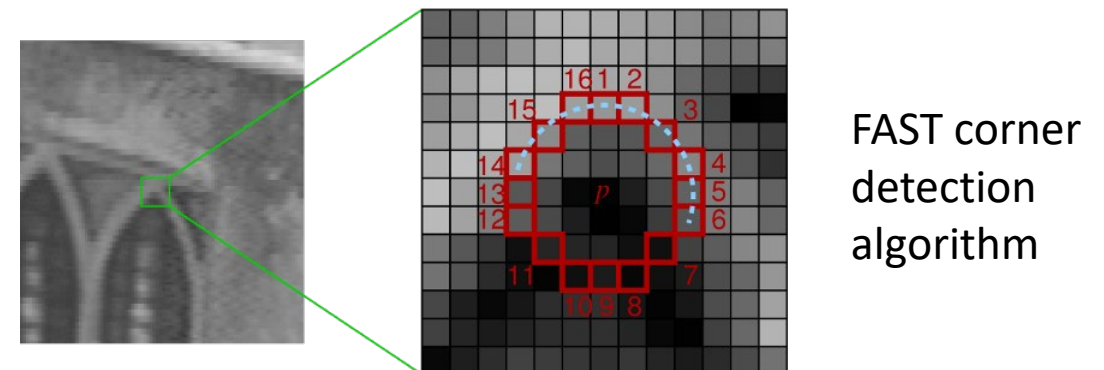
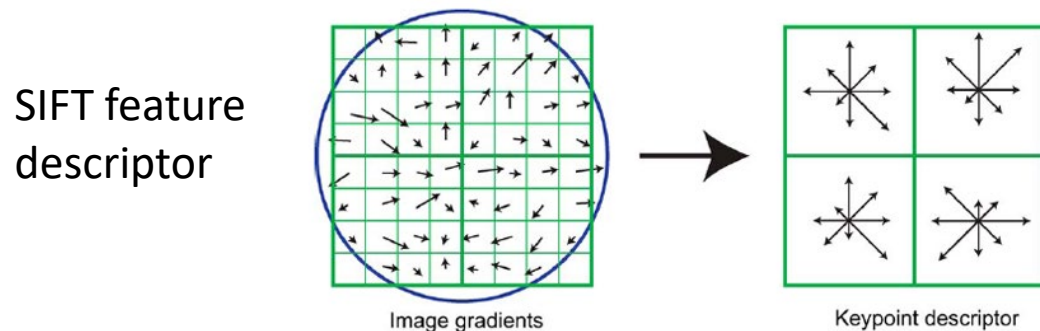


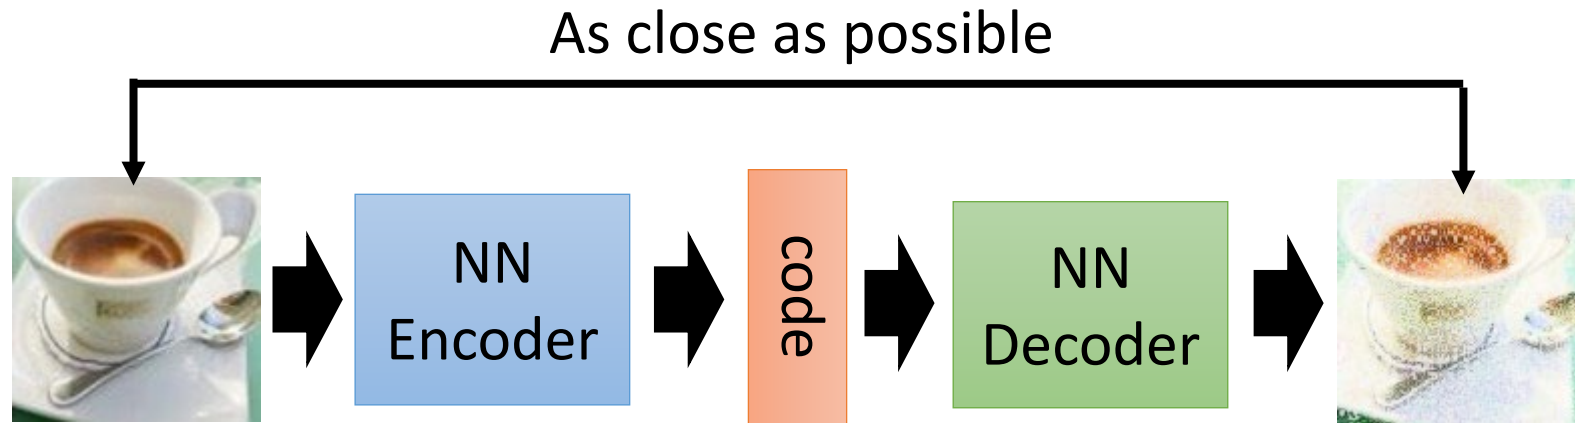
Image features (cont)

- What if we **learned the features to detect**?
- We need a system that can do Representation Learning (or Feature Learning).

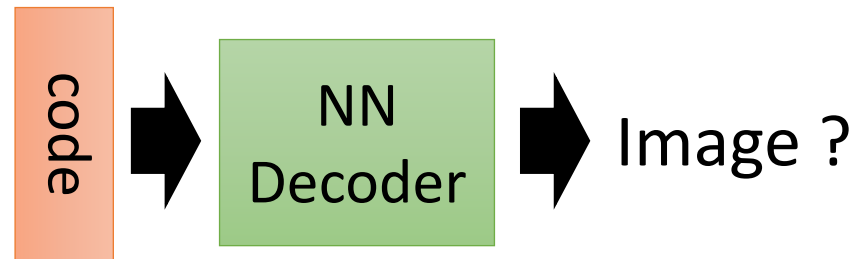
Representation Learning: technique that allows a system to automatically find relevant features for a given task. Replaces manual feature engineering.

- Multiple techniques for this:
 - Unsupervised (K-means, PCA, ...).
 - Supervised (Sup. Dictionary learning, **Neural Networks!**)

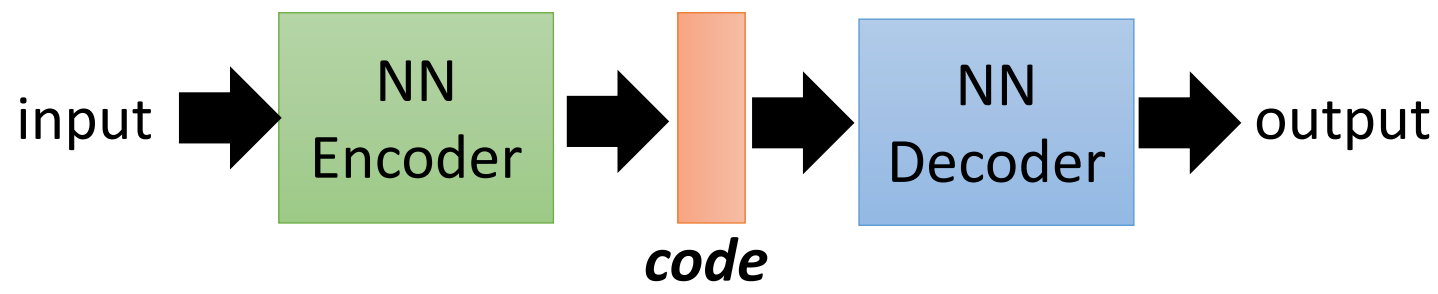
Auto-encoder



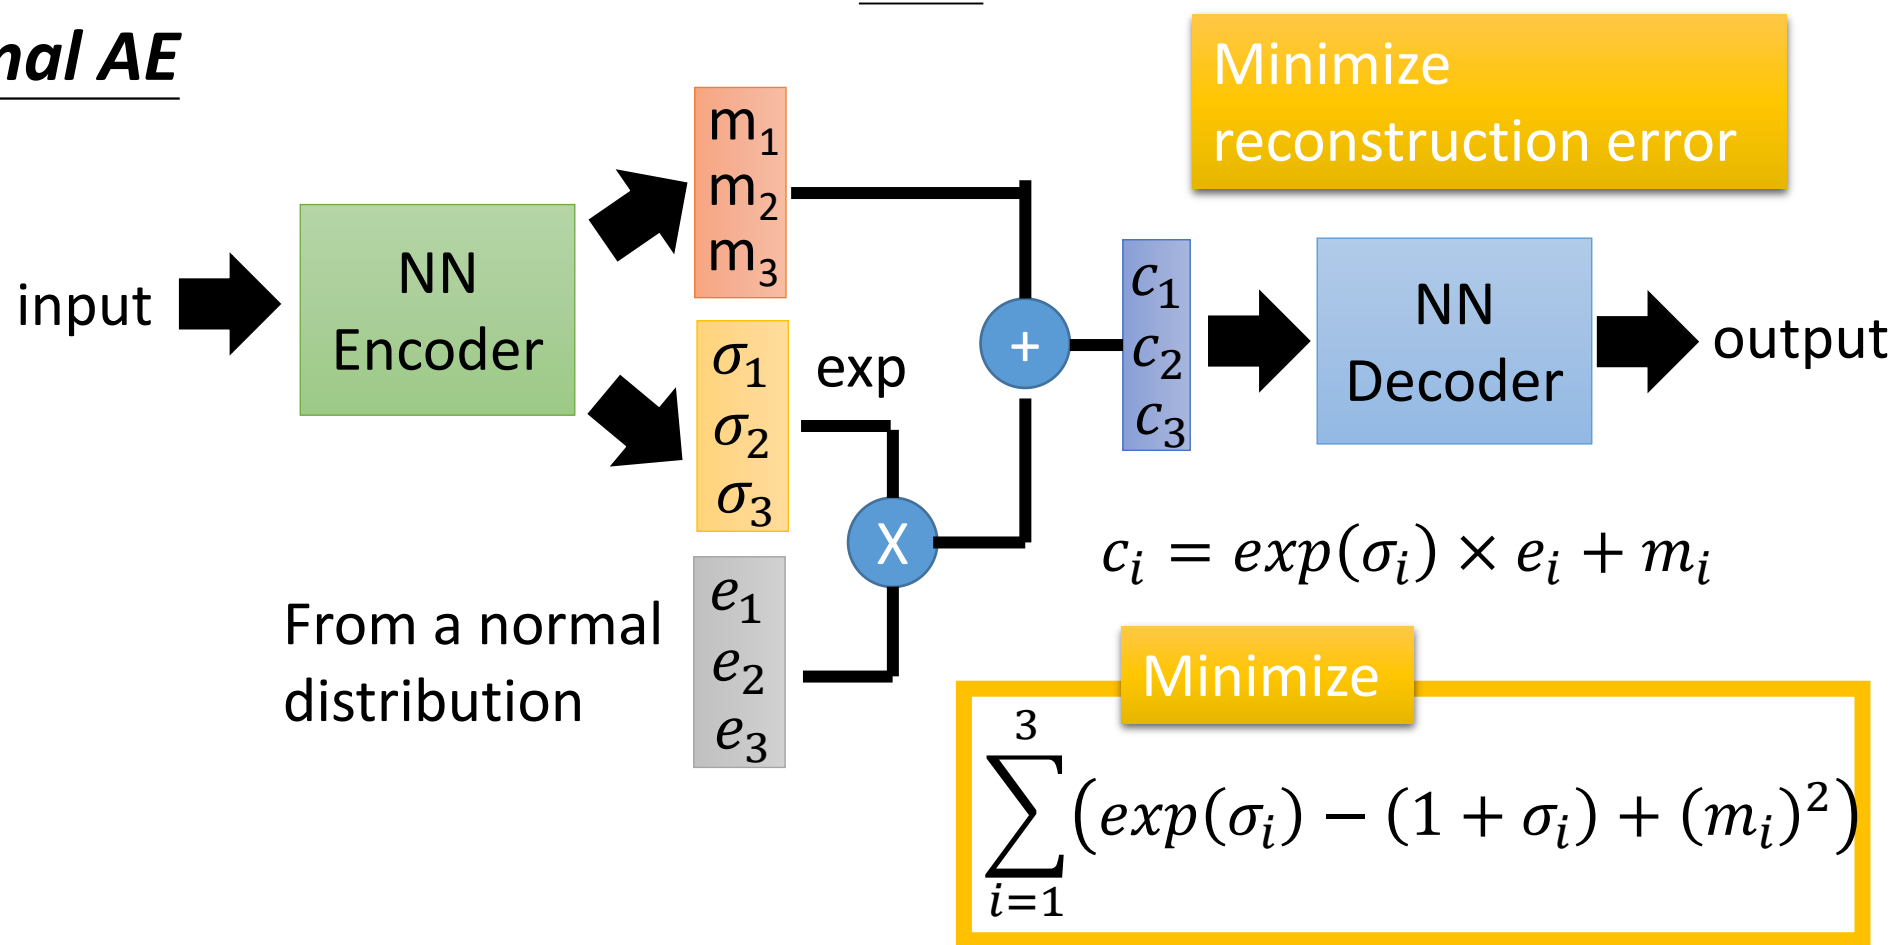
Randomly generate
a vector as code



Auto-encoder

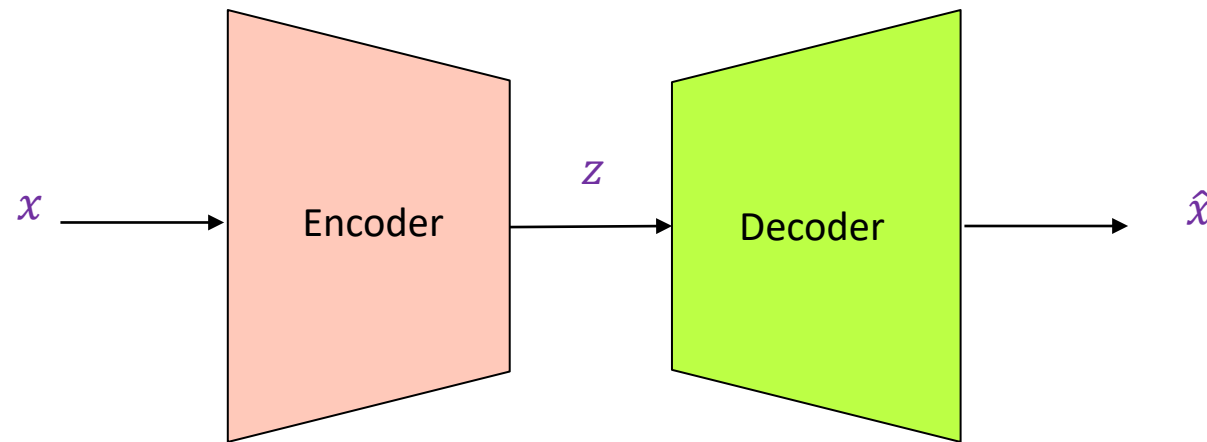


Variational AE



Variational autoencoders: Overview

- Probabilistic formulation based on *variational Bayes* framework
- At training time, jointly learn *encoder* and *decoder* by maximizing (a bound on) the data likelihood
- At test time, discard encoder and use decoder to sample from the learned distribution



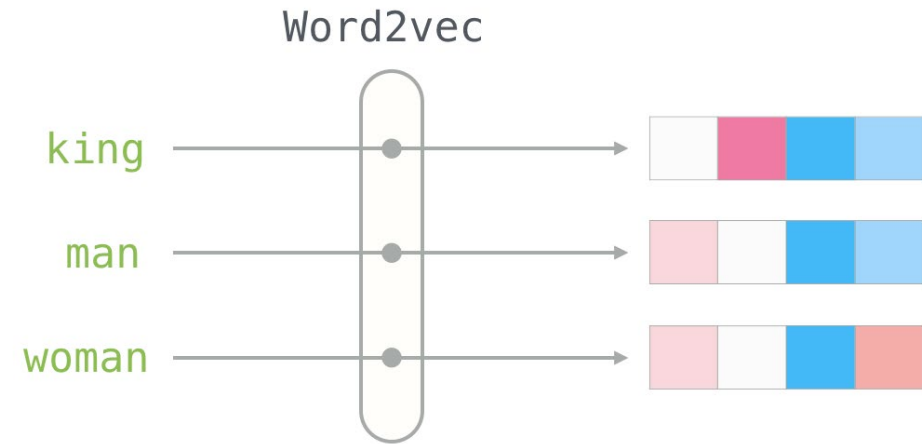
Embeddings

Neural network embeddings have 3 primary purposes:

- Finding nearest neighbors in the embedding space. These can be used to make recommendations based on user interests or cluster categories.
- As input to a machine learning model for a supervised task.
- For visualization of concepts and relations between categories.

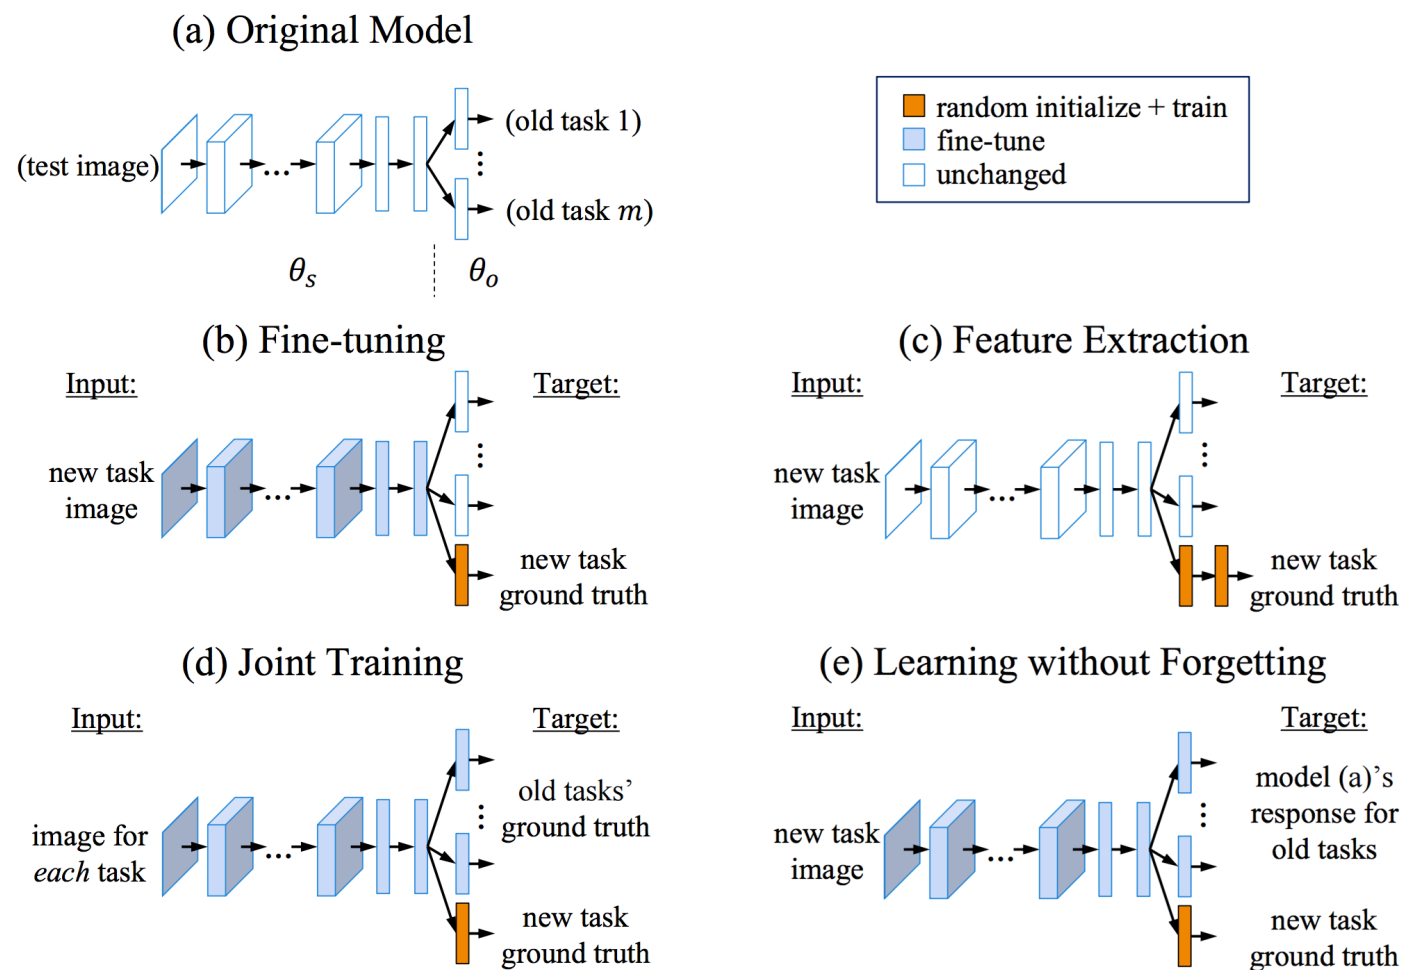
Word2Vec

Distributed vector representation for words



Mikolov, Tomas, Ilya Sutskever, Kai Chen, Greg S. Corrado, and Jeff Dean. "Distributed representations of words and phrases and their compositionality." In *Advances in neural information processing systems*, pp. 3111-3119. 2013.

Mixing Representations - LwF



Li & Hoiem
ECCV'16

Main drawbacks of ANNs (MLPs)

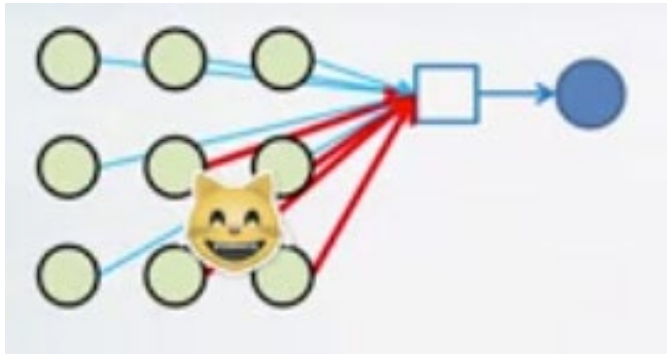
MLPs use one perceptron for each input (e.g. pixel in an image, multiplied by 3 in RGB case). The amount of weights **rapidly becomes unmanageable** for large images.

Training difficulties arise, overfitting can appear.

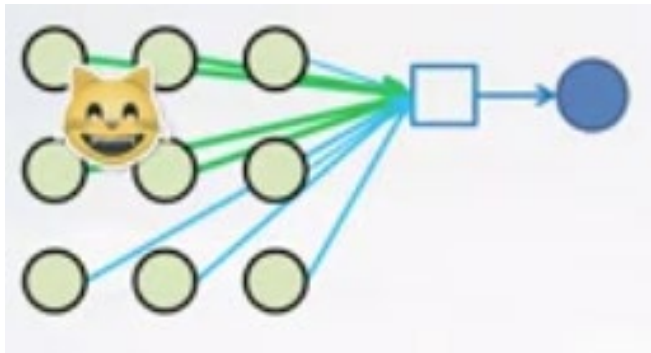
MLPs react differently to an input (images) and its shifted version – **they are not translation invariant**.

Drawbacks of ANNs and MLPs

Imagine we want to build a cat detector with an MLP.



In this case, the **red weights** will be modified to better recognize cats



In this case, the **green weights** will be modified.

We are learning redundant features. Approach is not robust, as cats could appear in yet another position.

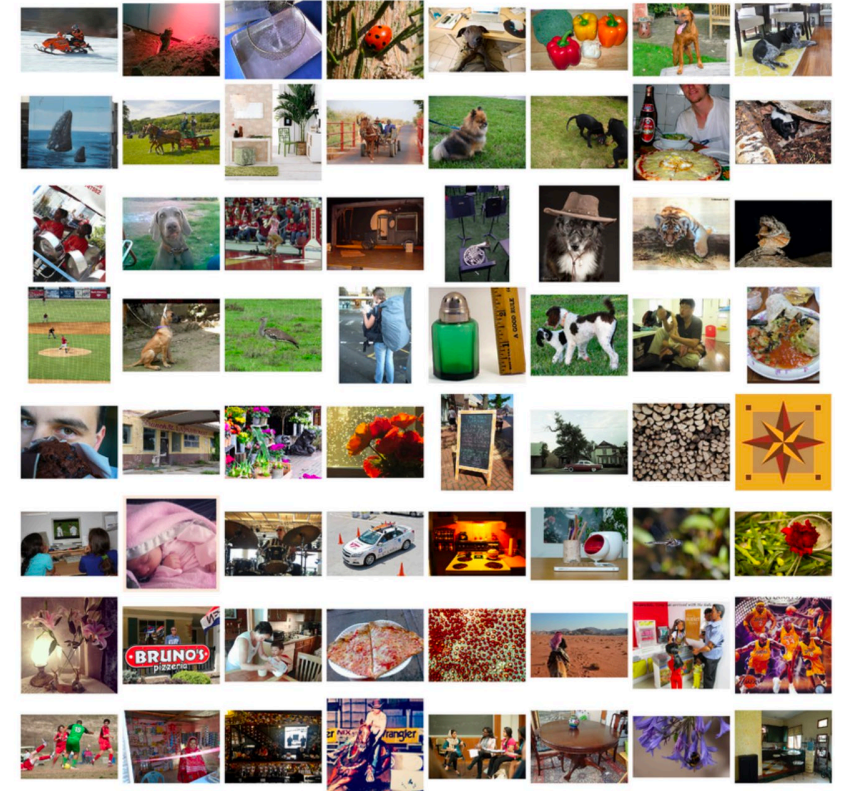
Drawbacks

- Example: CIFAR10
- Simple 32x32 color images (3 channels)
- Each pixel is a feature: an MLP would have $32 \times 32 \times 3 + 1 = 3073$ weights per neuron!



Drawbacks

- Example: ImageNet
- Images are usually $224 \times 224 \times 3$: an MLP would have **150129 weights per neuron**. If the first layer of the MLP is around 128 nodes, which is small, this already becomes very heavy to calculate.
- Model complexity is extremely high: overfitting.



Images are Local and Hierarchical



Nearby pixels are more strongly related than distant ones.

Objects are built up out of smaller parts.

Images are Invariant



Basics of CNNs

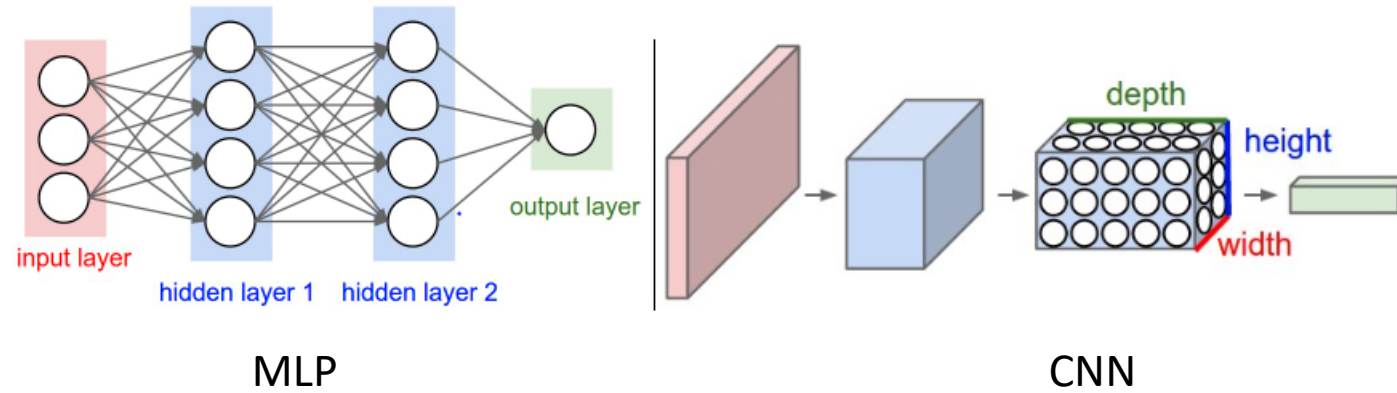
We know that MLPs:

- Do not scale well for images
- Ignore the information brought by pixel position and correlation with neighbors
- Cannot handle translations

The general idea of CNNs is to intelligently adapt to properties of images:

- Pixel position and neighborhood have semantic meanings.
- Elements of interest can appear anywhere in the image.

Basics of CNNs

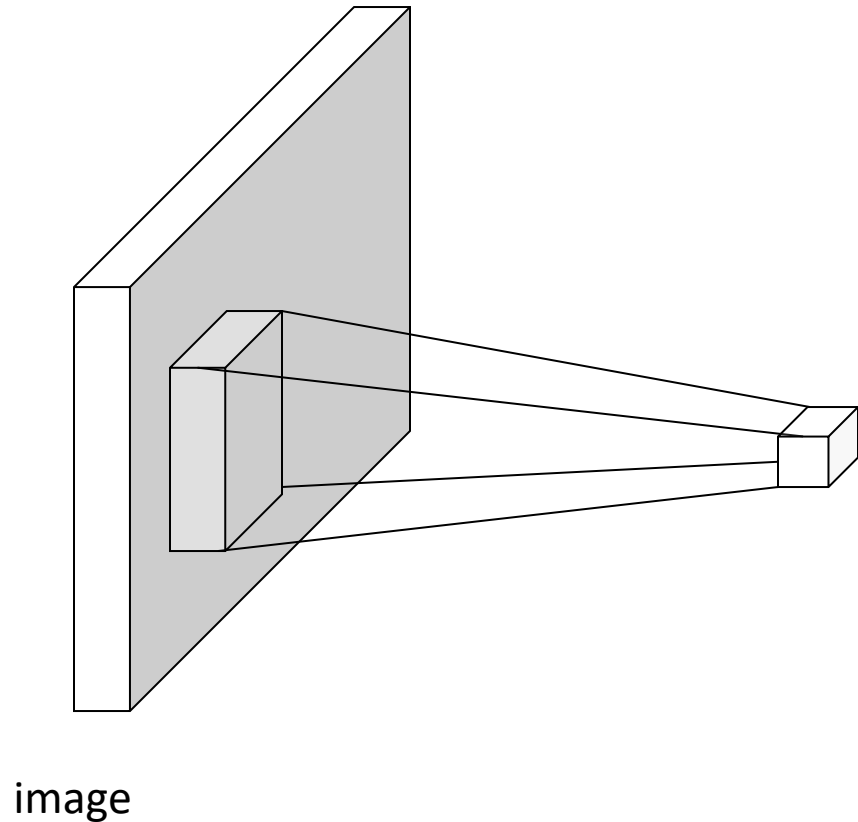


CNNs are also composed of layers, but those layers are not fully connected: they have **filters**, sets of cube-shaped weights that are applied throughout the image. Each 2D slice of the filters are called **kernels**.

These filters introduce **translation invariance** and **parameter sharing**.

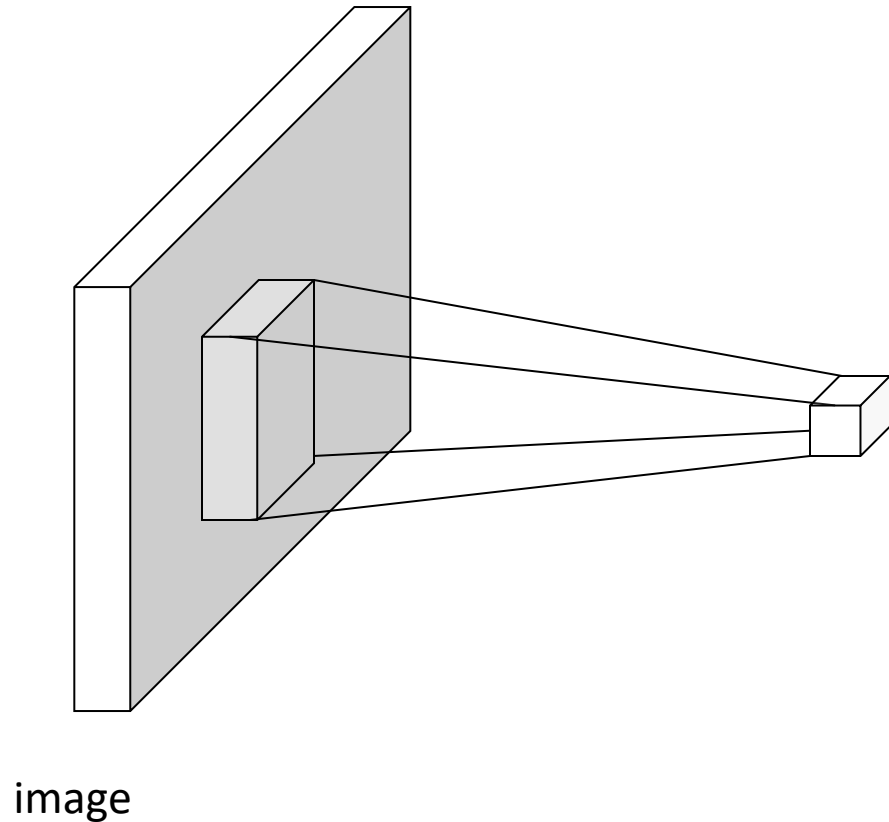
How are they applied? Convolutions!

Convolutional architecture



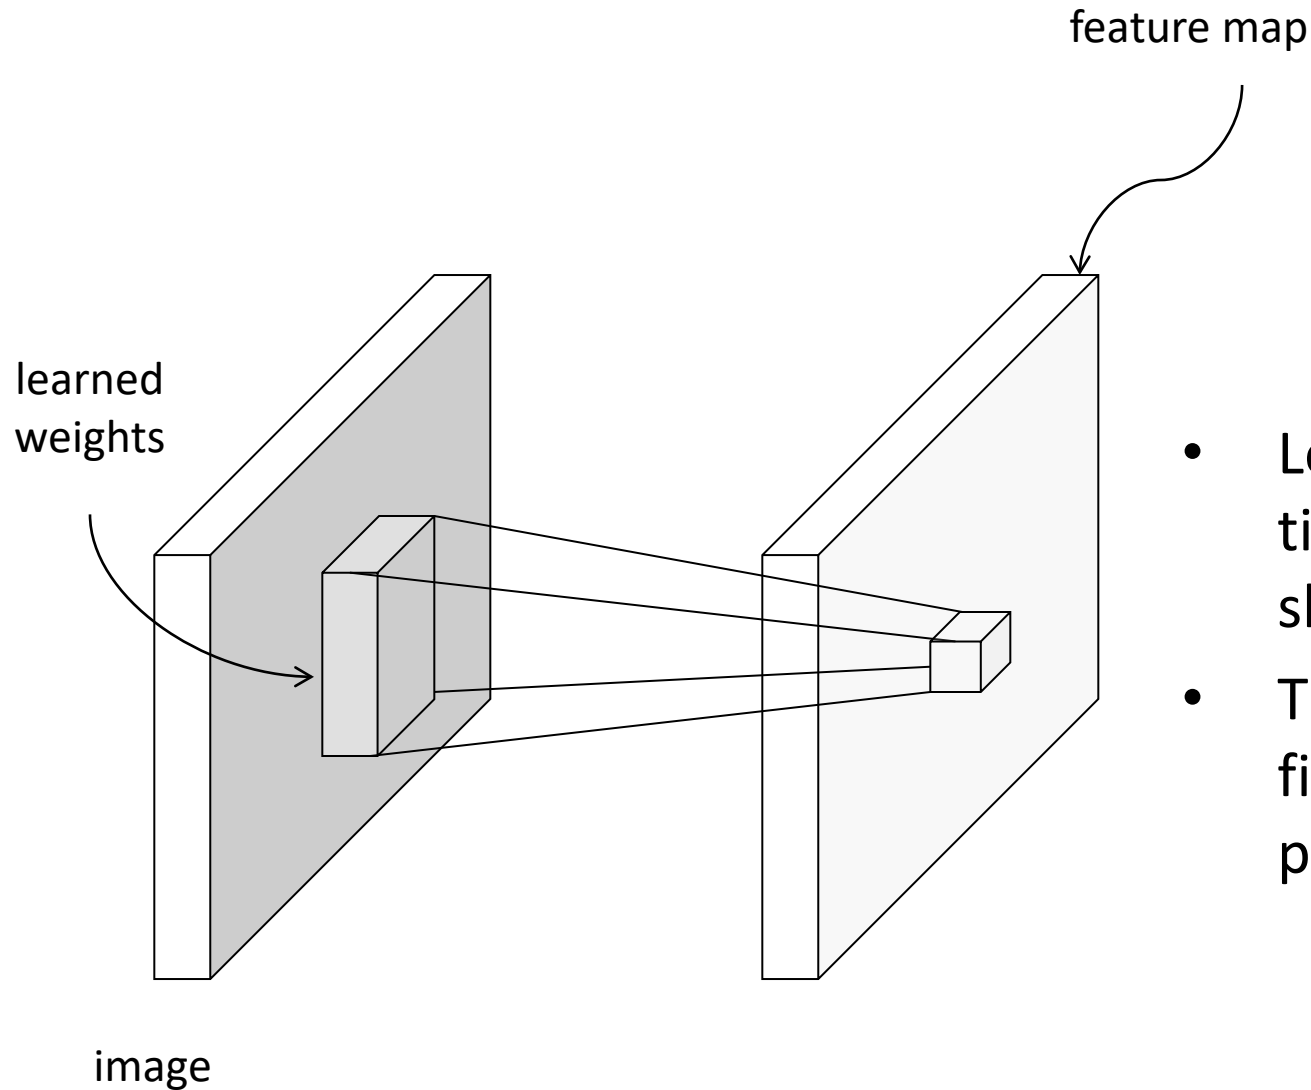
- Let's limit the *receptive fields* of units, tile them over the input image, and share their weights

Convolutional architecture



- Let's limit the *receptive fields* of units, tile them over the input image, and share their weights

Convolutional architecture



- Let's limit the *receptive fields* of units, tile them over the input image, and share their weights
- This is equivalent to sliding the learned filter over the image, computing dot products at every location

Convolution example

Input

x_{11}	x_{12}	x_{13}	x_{14}	x_{15}	x_{16}
x_{21}	x_{22}	x_{23}	x_{24}	x_{25}	x_{26}
x_{31}	x_{32}	x_{33}	x_{34}	x_{35}	x_{36}
x_{41}	x_{42}	x_{43}	x_{44}	x_{45}	x_{46}
x_{51}	x_{52}	x_{53}	x_{54}	x_{55}	x_{56}

Filter

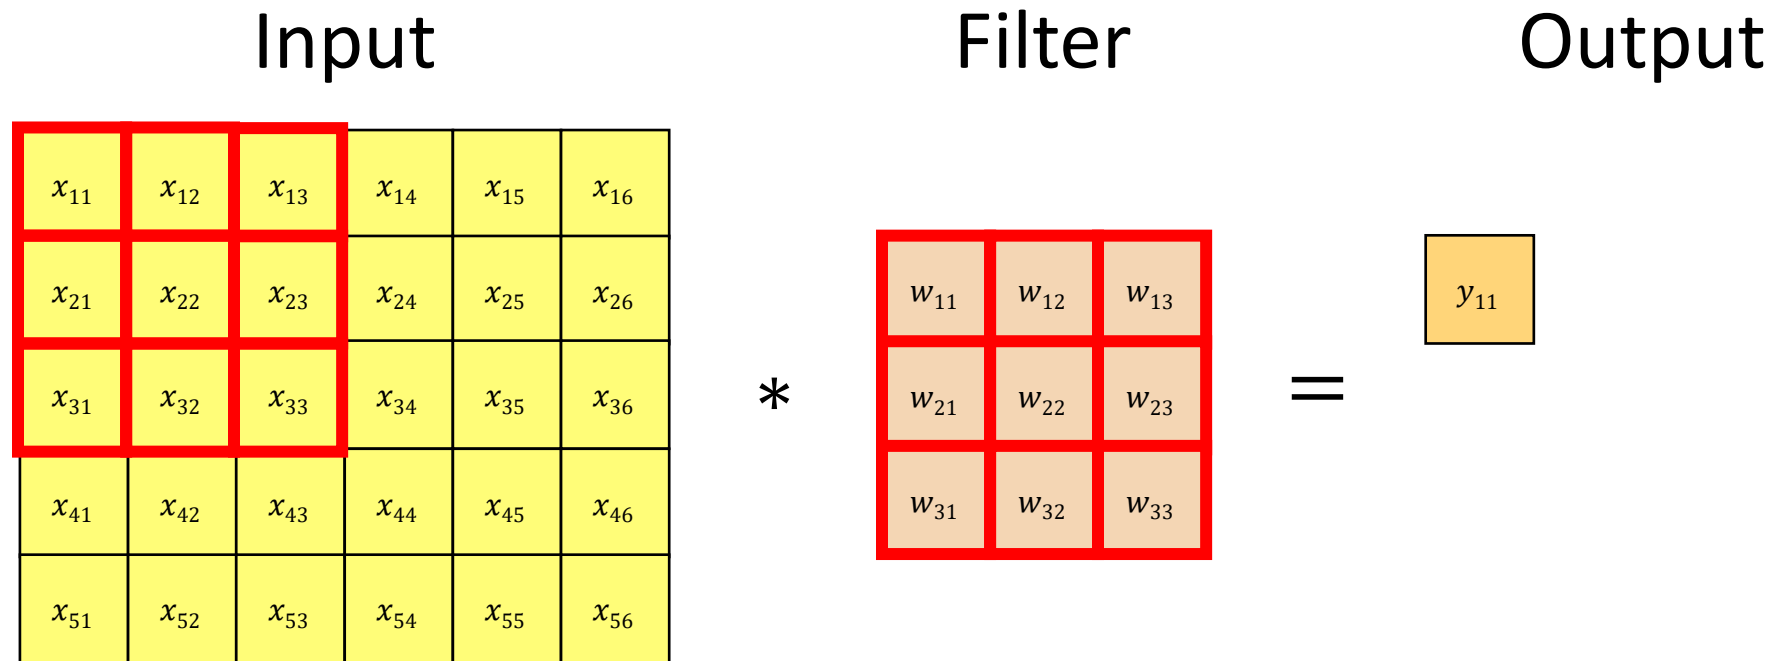
w_{11}	w_{12}	w_{13}
w_{21}	w_{22}	w_{23}
w_{31}	w_{32}	w_{33}

*

=

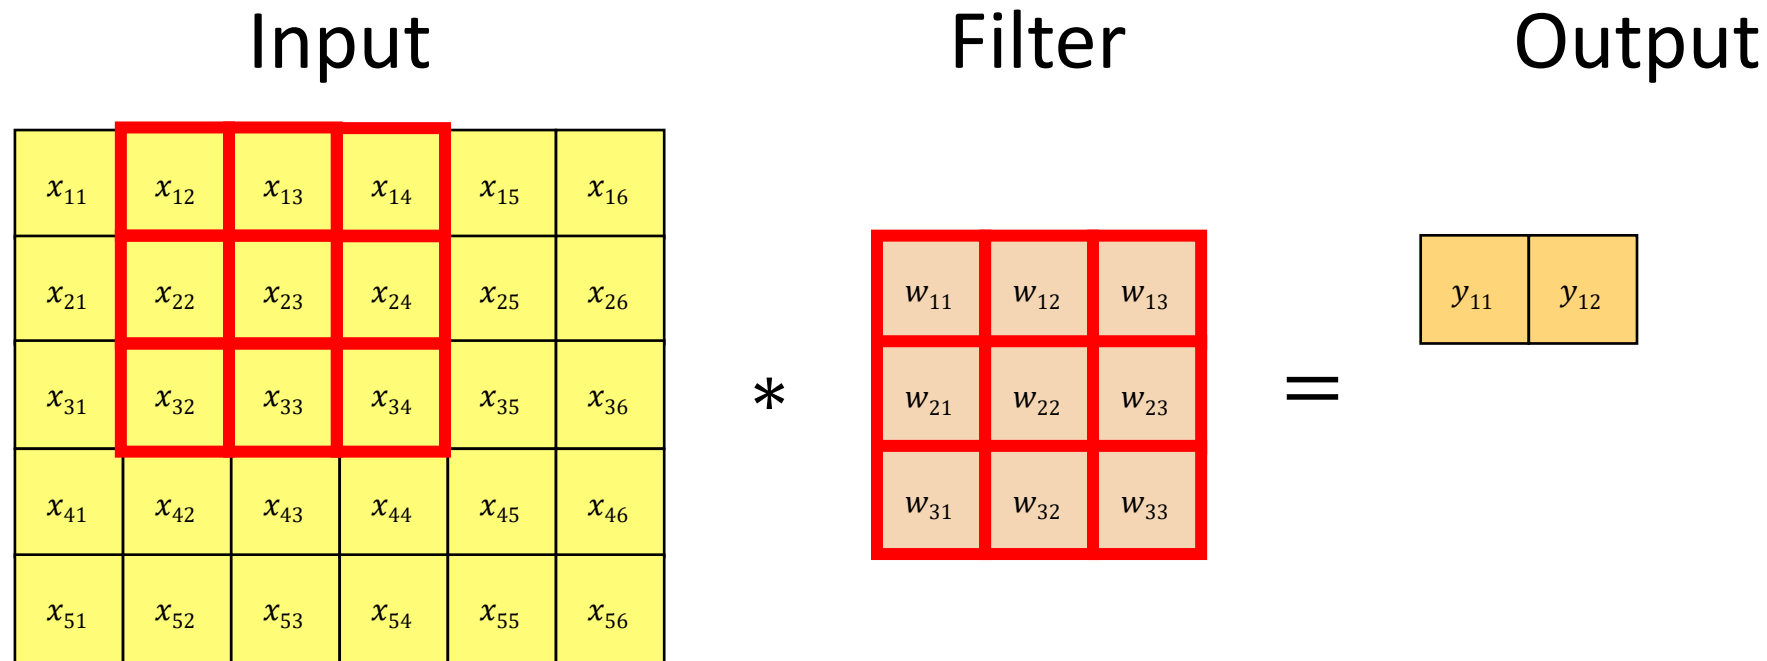
Output

Convolution example



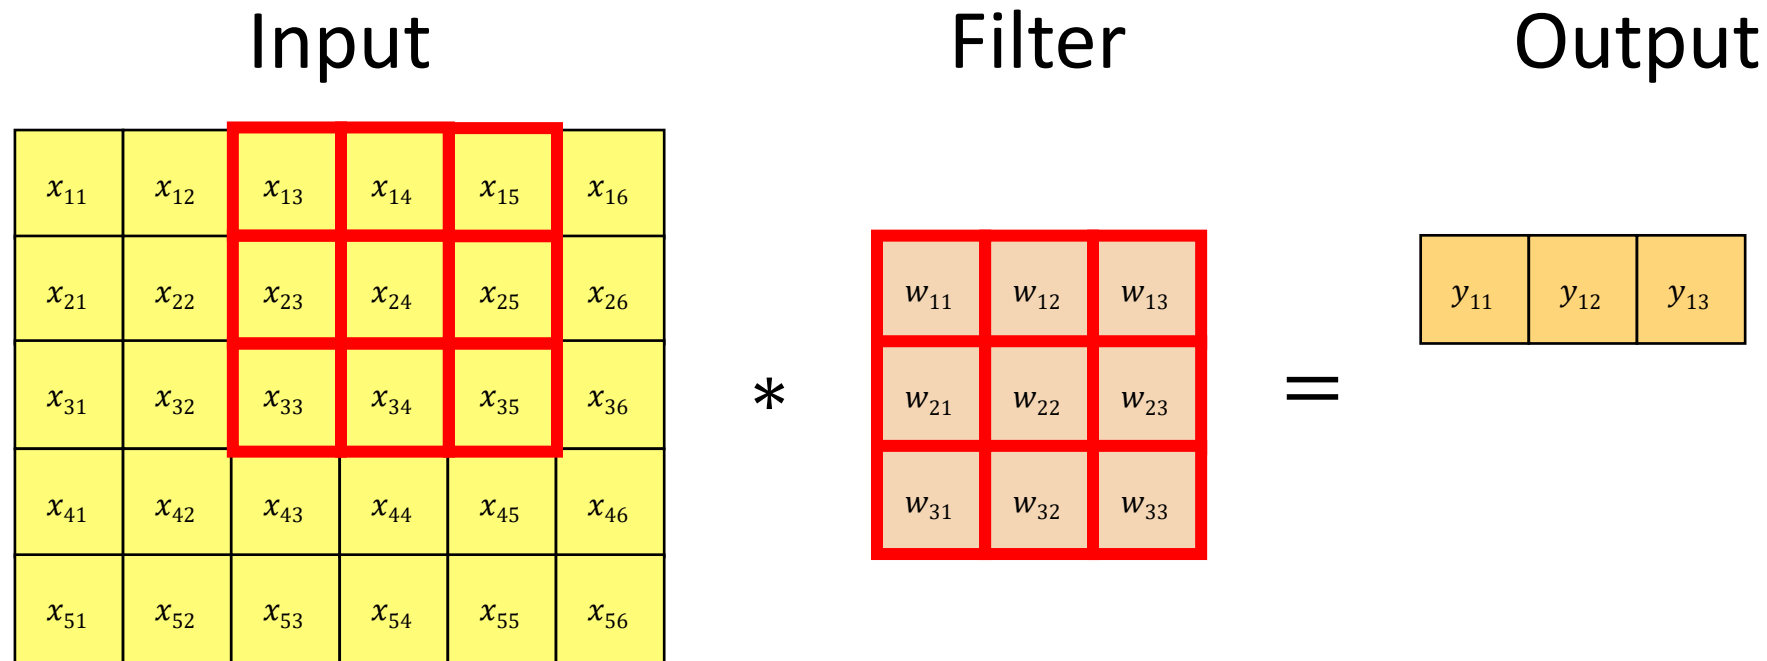
$$y_{11} = x_{11} \cdot w_{11} + x_{12} \cdot w_{12} + x_{13} \cdot w_{13} + \dots + x_{33} \cdot w_{33}$$

Convolution example



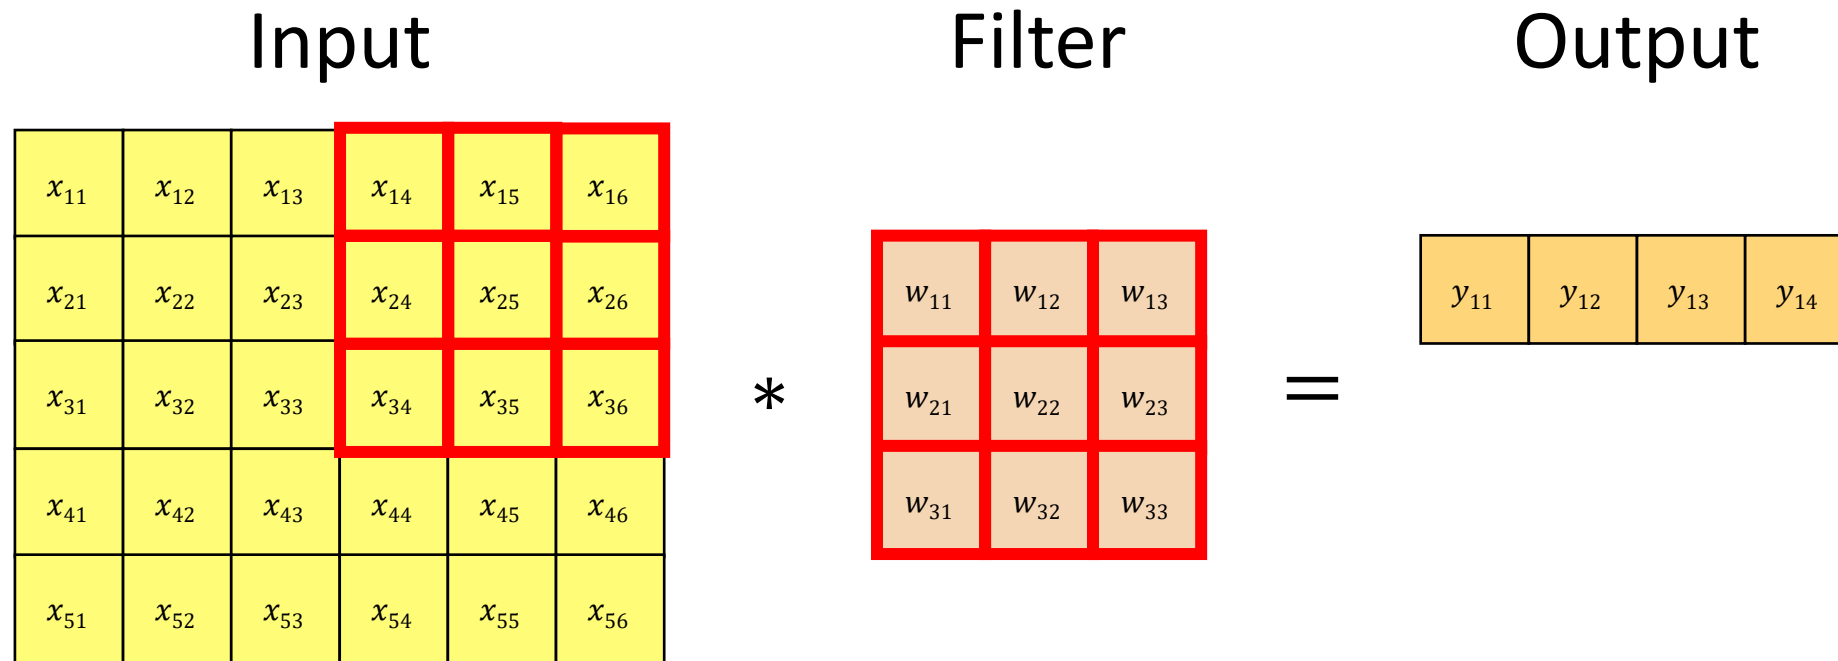
$$y_{12} = x_{12} \cdot w_{11} + x_{13} \cdot w_{12} + x_{14} \cdot w_{13} + \dots + x_{34} \cdot w_{33}$$

Convolution example



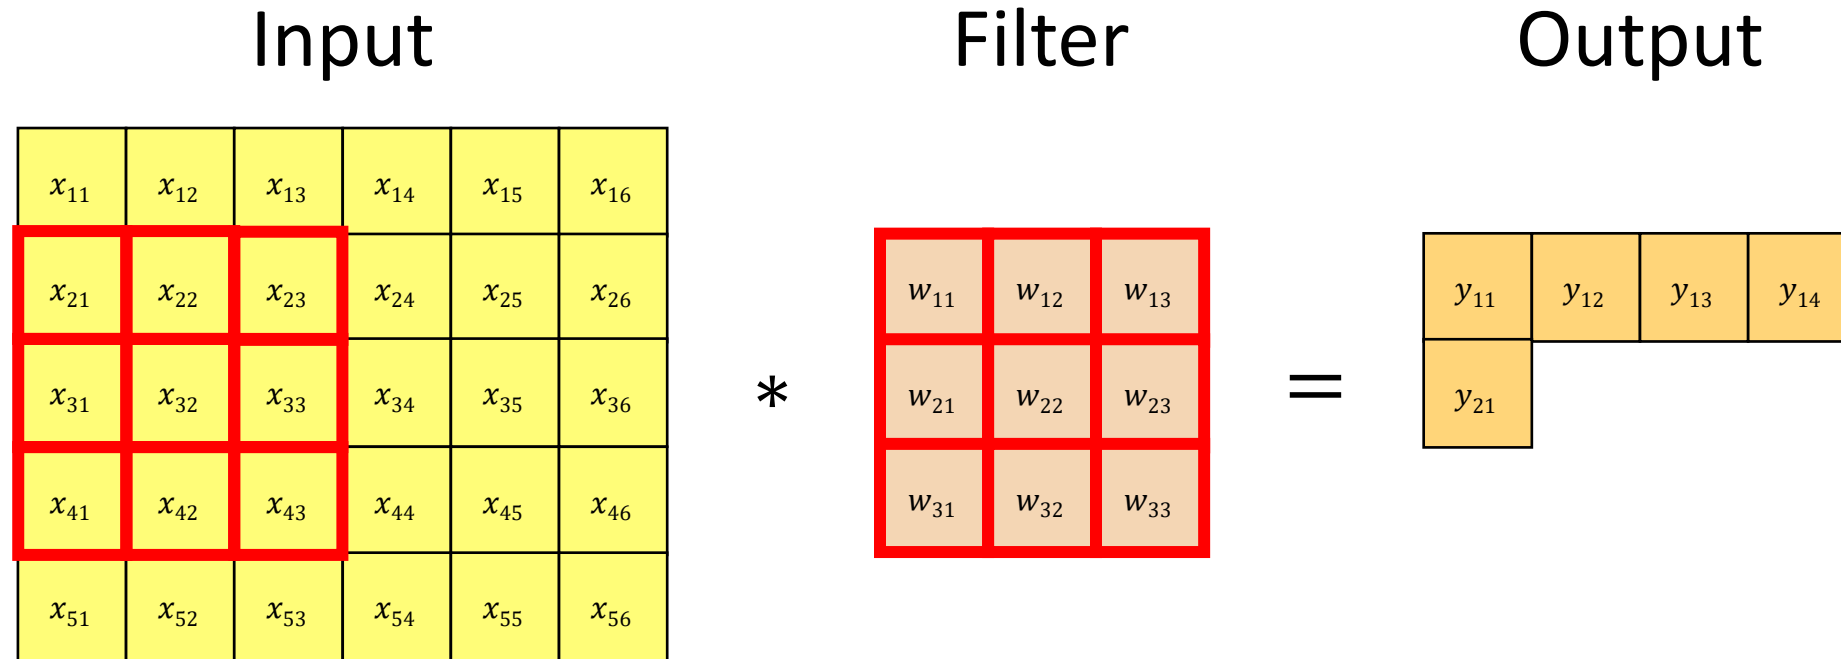
$$y_{13} = x_{13} \cdot w_{11} + x_{14} \cdot w_{12} + x_{15} \cdot w_{13} + \dots + x_{35} \cdot w_{33}$$

Convolution example



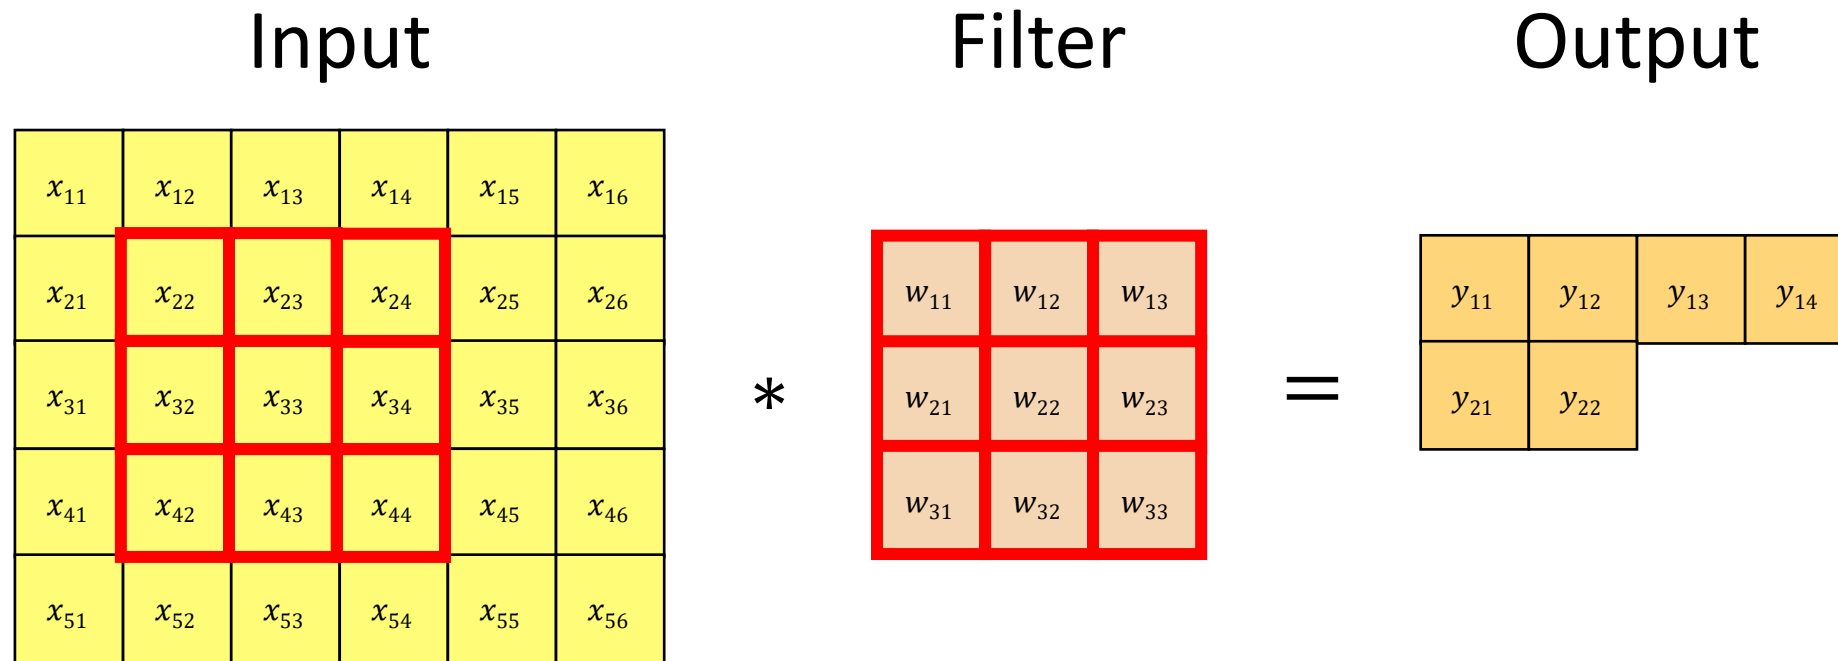
$$y_{14} = x_{14} \cdot w_{11} + x_{15} \cdot w_{12} + x_{16} \cdot w_{13} + \dots + x_{36} \cdot w_{33}$$

Convolution example



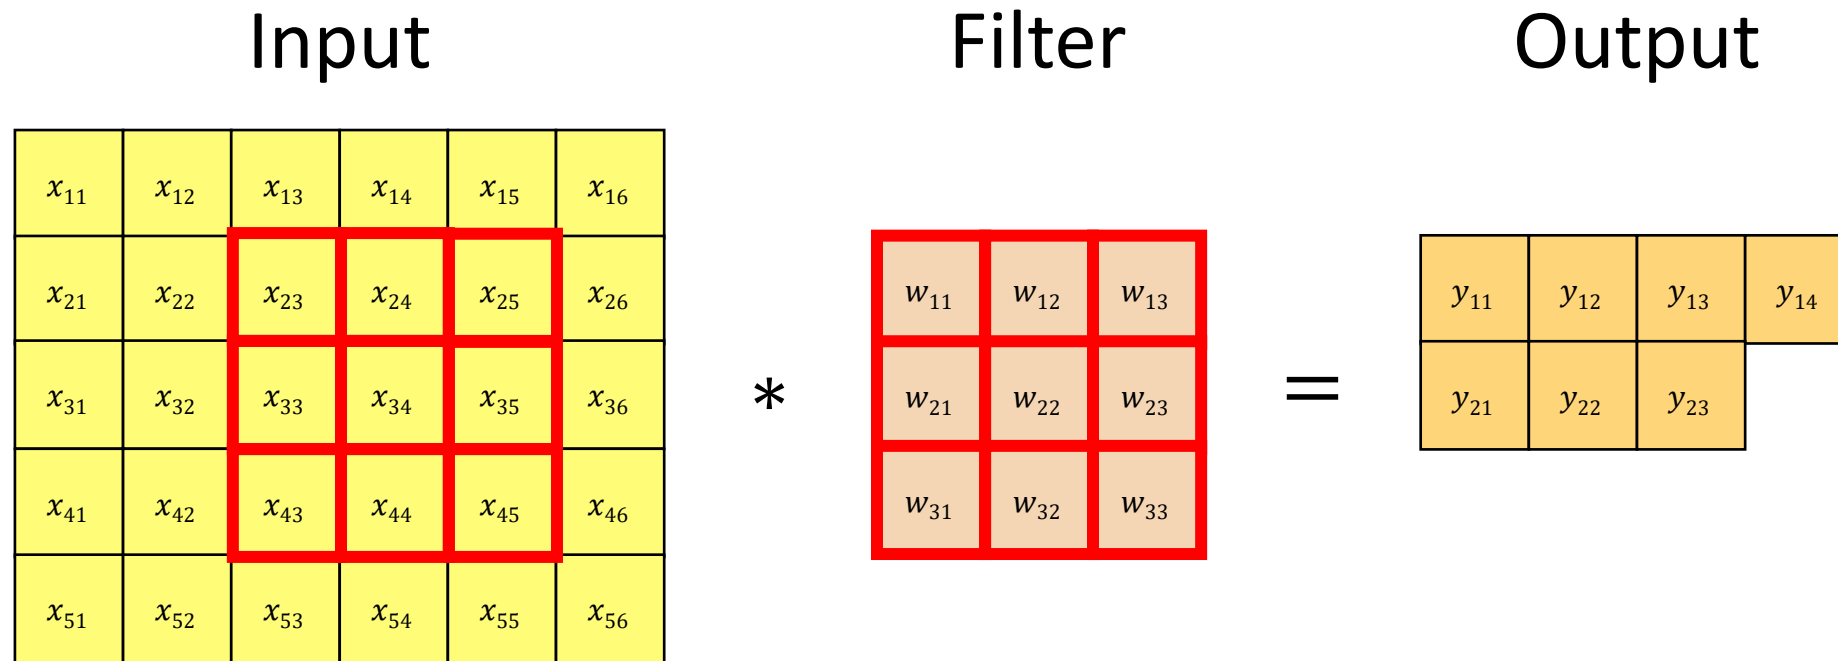
$$y_{21} = x_{21} \cdot w_{11} + x_{22} \cdot w_{12} + x_{23} \cdot w_{13} + \dots + x_{43} \cdot w_{33}$$

Convolution example



$$y_{22} = x_{22} \cdot w_{11} + x_{23} \cdot w_{12} + x_{24} \cdot w_{13} + \dots + x_{44} \cdot w_{33}$$

Convolution example



$$y_{23} = x_{23} \cdot w_{11} + x_{24} \cdot w_{12} + x_{25} \cdot w_{13} + \dots + x_{45} \cdot w_{33}$$

Convolution example

Input

x_{11}	x_{12}	x_{13}	x_{14}	x_{15}	x_{16}
x_{21}	x_{22}	x_{23}	x_{24}	x_{25}	x_{26}
x_{31}	x_{32}	x_{33}	x_{34}	x_{35}	x_{36}
x_{41}	x_{42}	x_{43}	x_{44}	x_{45}	x_{46}
x_{51}	x_{52}	x_{53}	x_{54}	x_{55}	x_{56}

Filter

w_{11}	w_{12}	w_{13}
w_{21}	w_{22}	w_{23}
w_{31}	w_{32}	w_{33}

*

=

Output

y_{11}	y_{12}	y_{13}	y_{14}
y_{21}	y_{22}	y_{23}	y_{24}
y_{31}	y_{32}	y_{33}	y_{34}

Convolution and cross-correlation

- A **convolution** of f and g ($f * g$) is defined as the integral of the product, having one of the functions inverted and shifted:

$$(f * g)(t) = \int_a f(a)g(t - a)da$$

Function is
inverted and
shifted left by t

- Discrete convolution:

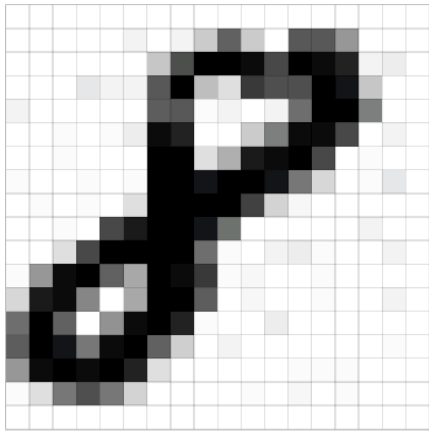
$$(f * g)(t) = \sum_{a=-\infty}^{\infty} f(a)g(t - a)$$

- Discrete cross-correlation:



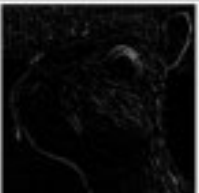


$$(f \star g)(t) = \sum_{a=-\infty}^{\infty} f(a)g(t + a)$$

Why does this make sense?

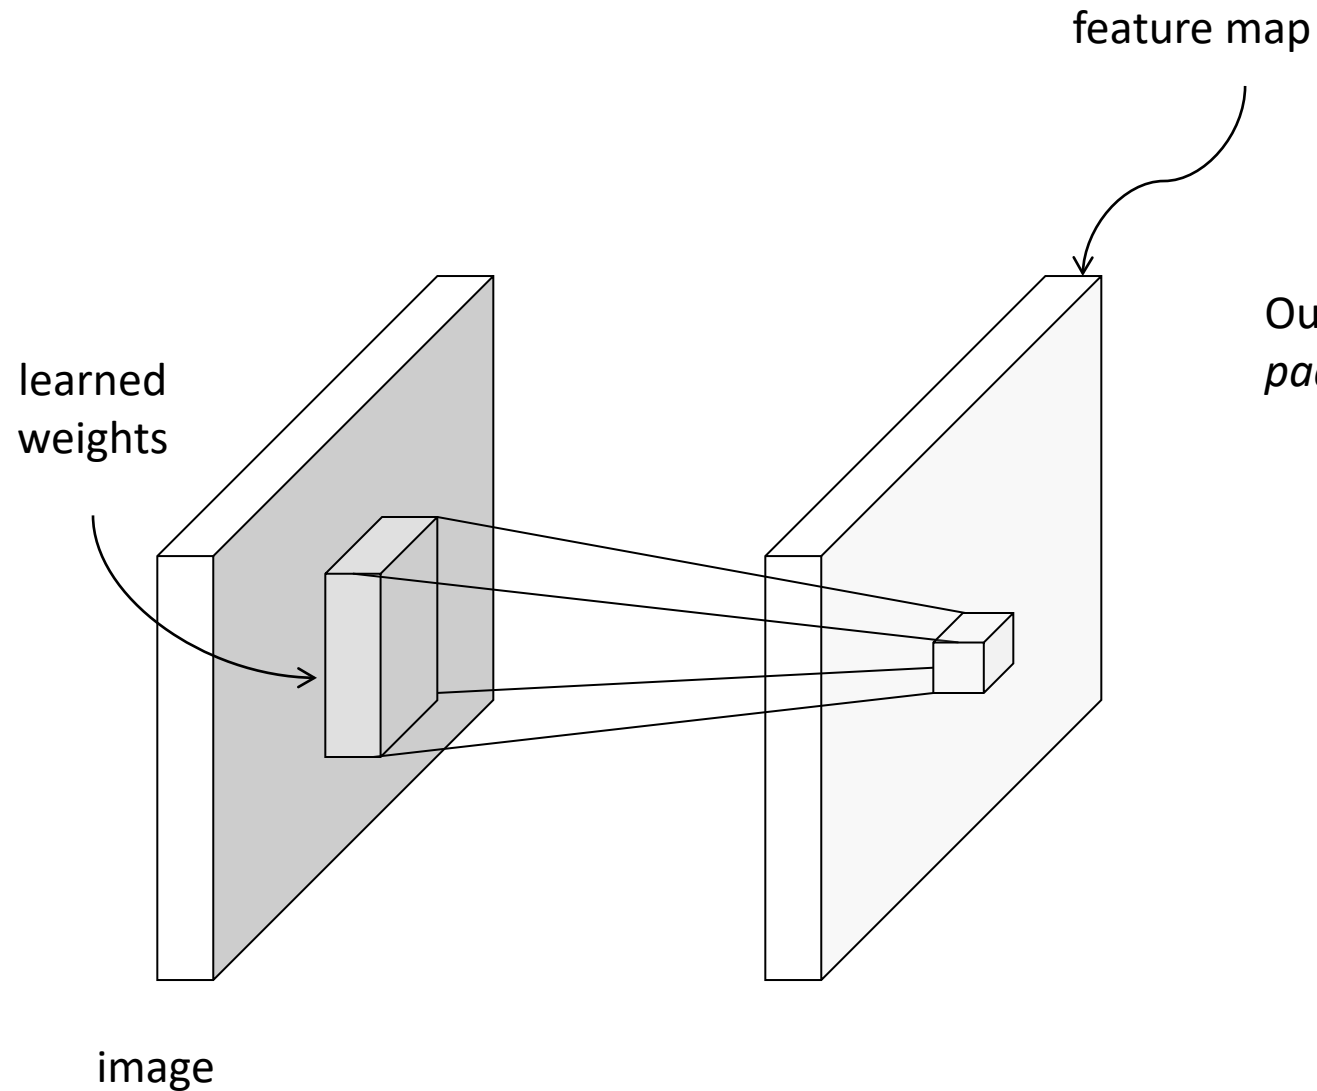
- In image is just a matrix of pixels.



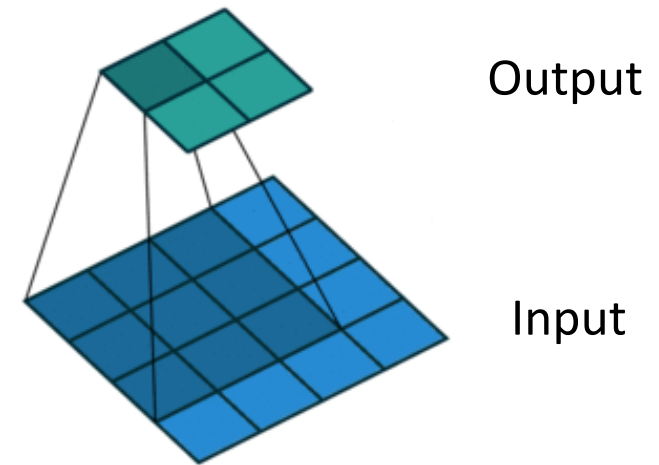
- Convoluting the image with a filter produces a feature map that highlights the presence of a given feature in the image.

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	

Convolutional architecture

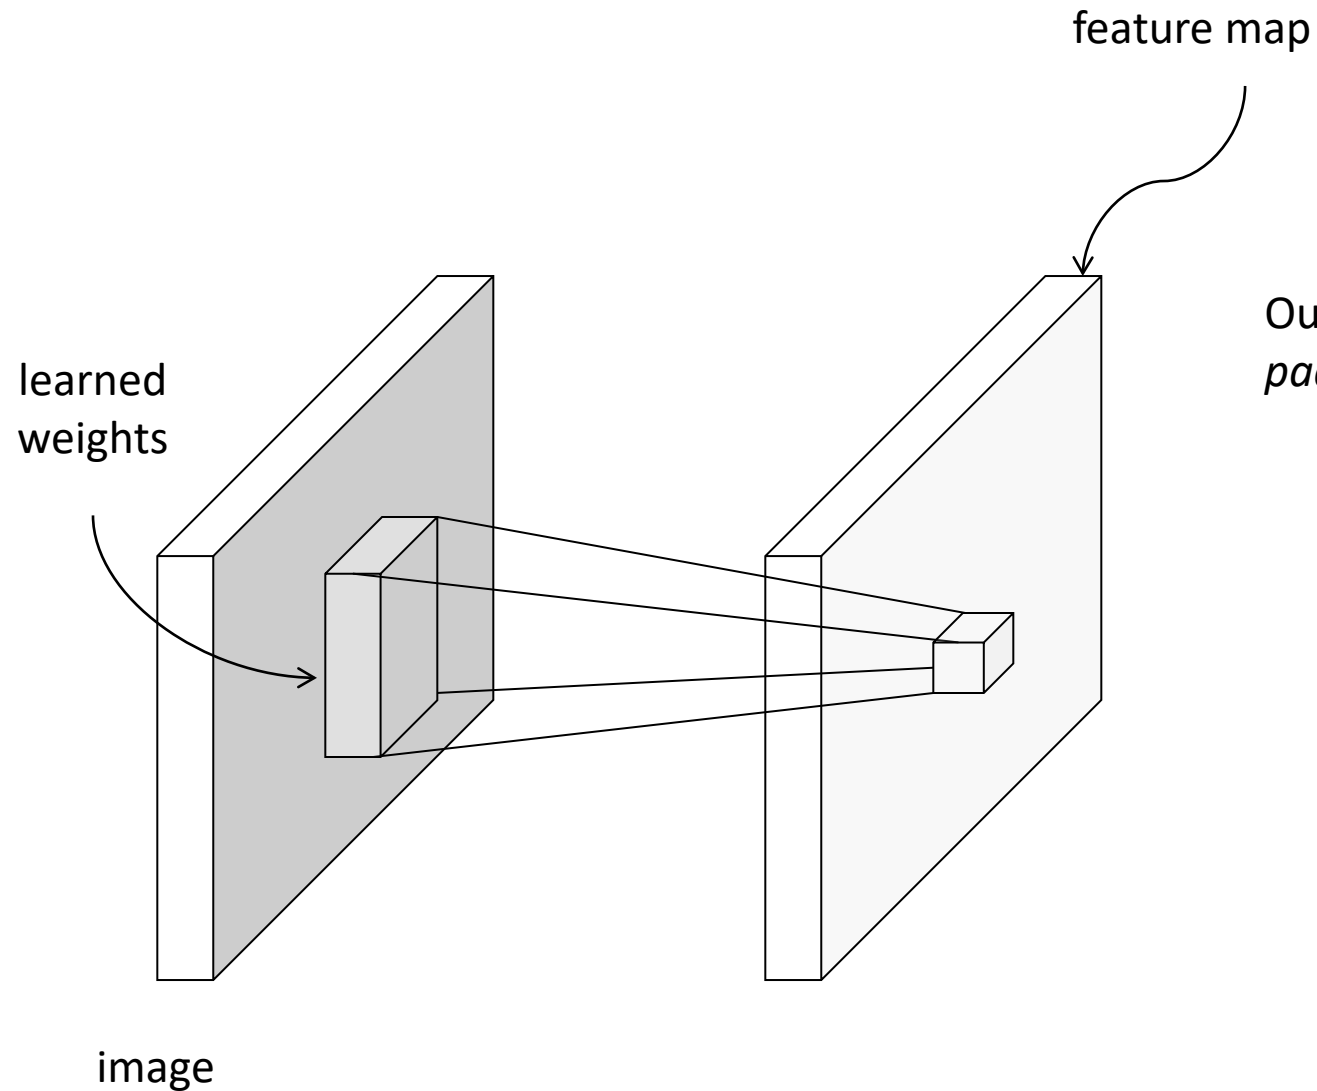


Output feature map resolution depends on *padding* and *stride*

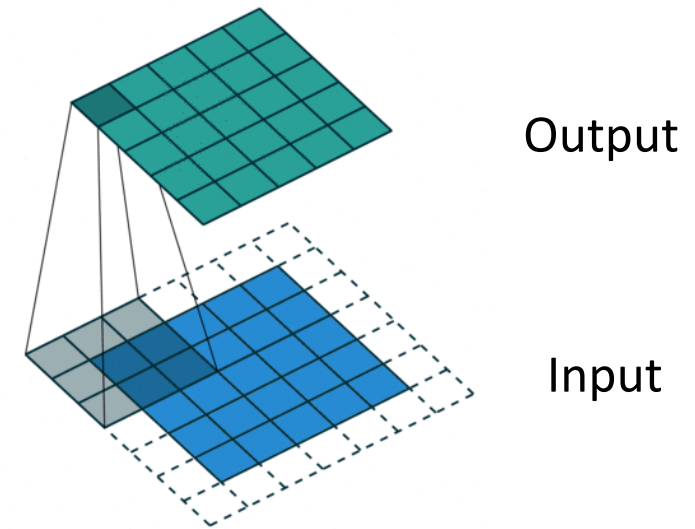


No padding, stride 1

Convolutional architecture

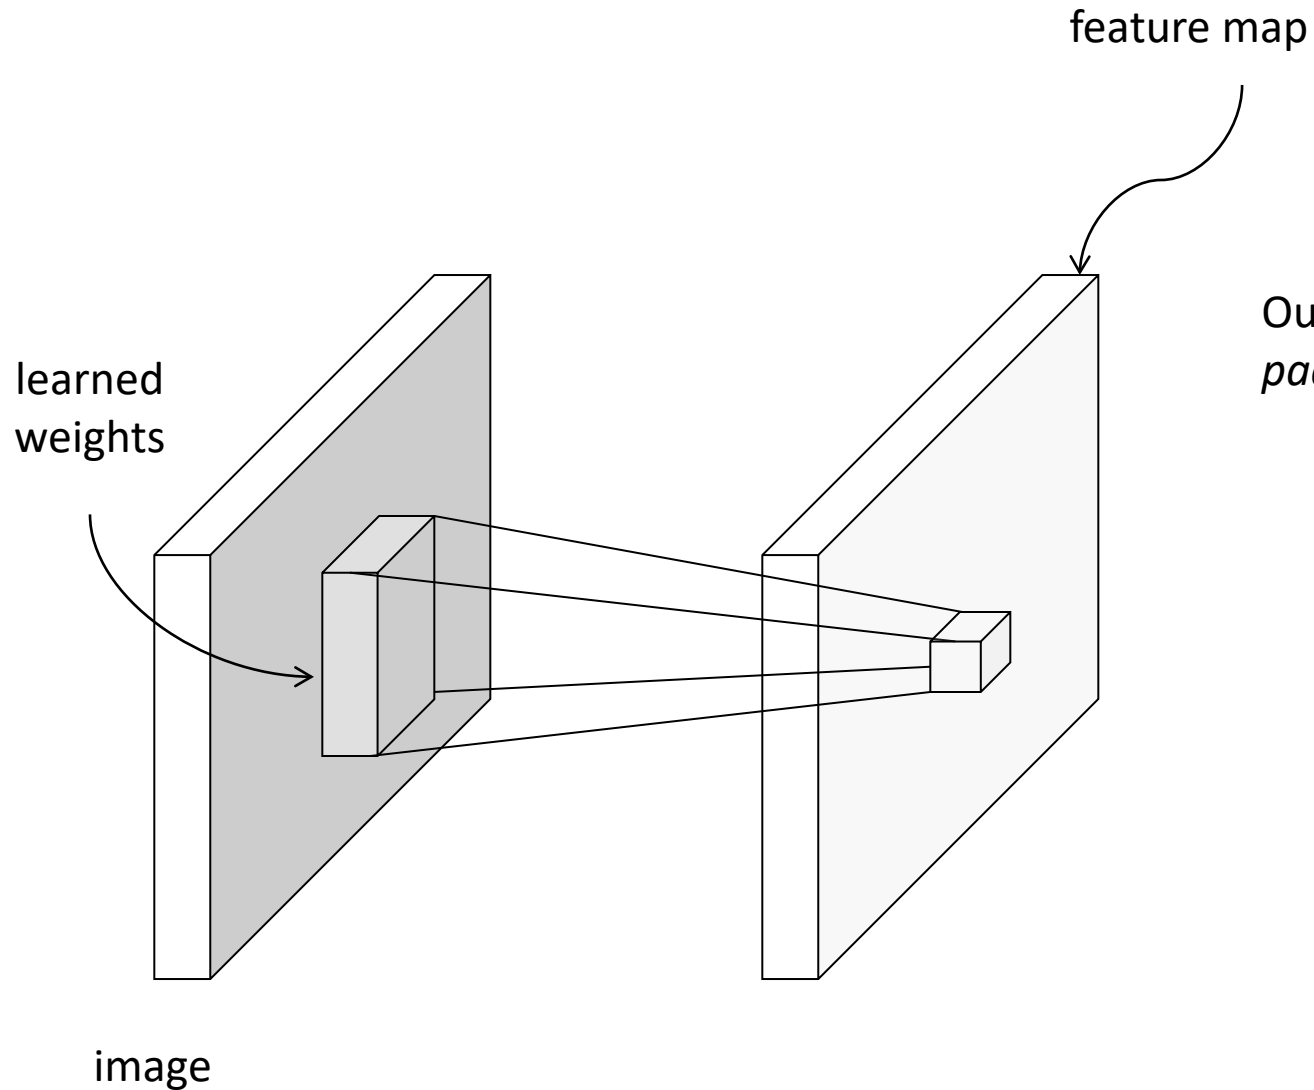


Output feature map resolution depends on *padding* and *stride*

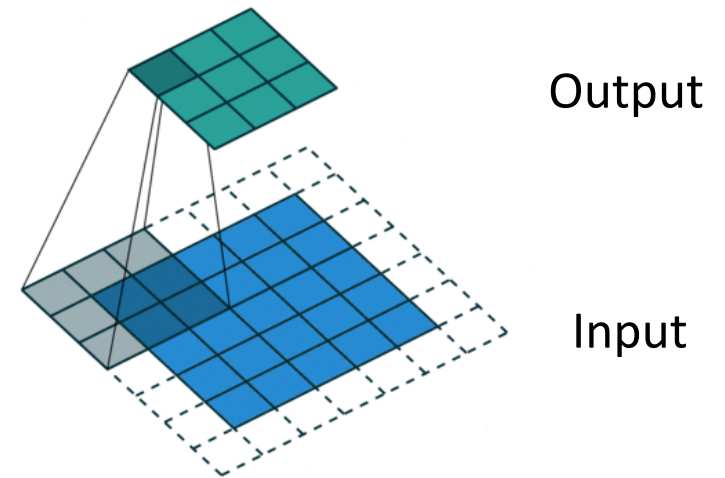


With padding, stride 1

Convolutional architecture

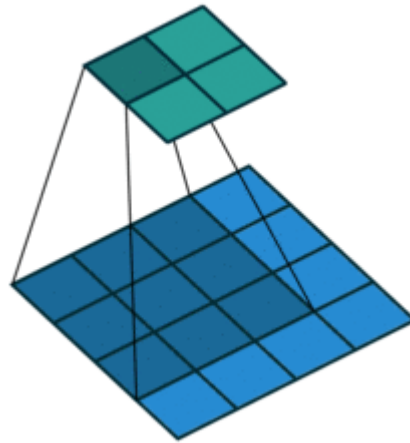


Output feature map resolution depends on *padding* and *stride*



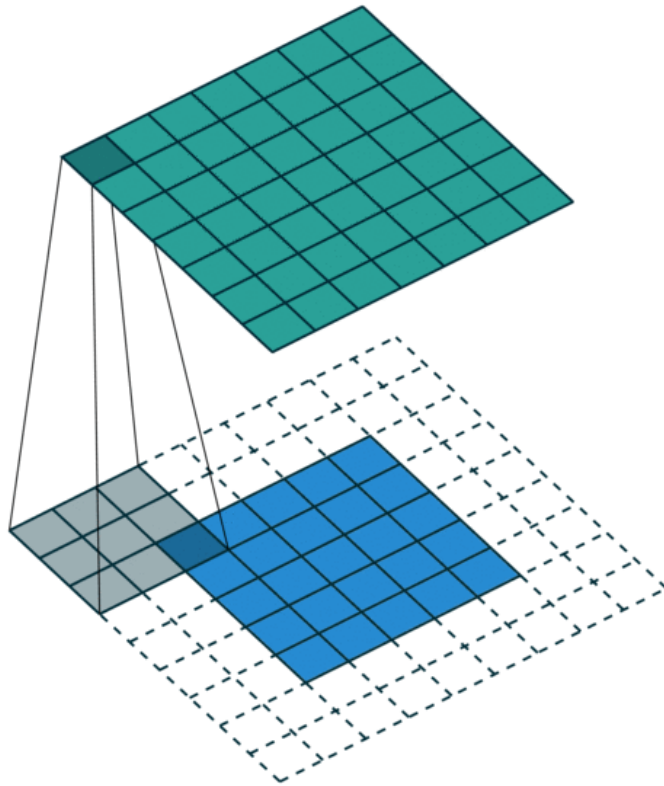
Convolutions – what happens at the edges?

If we apply convolutions on a normal image, the result will be down-sampled by an amount depending on the size of the filter.

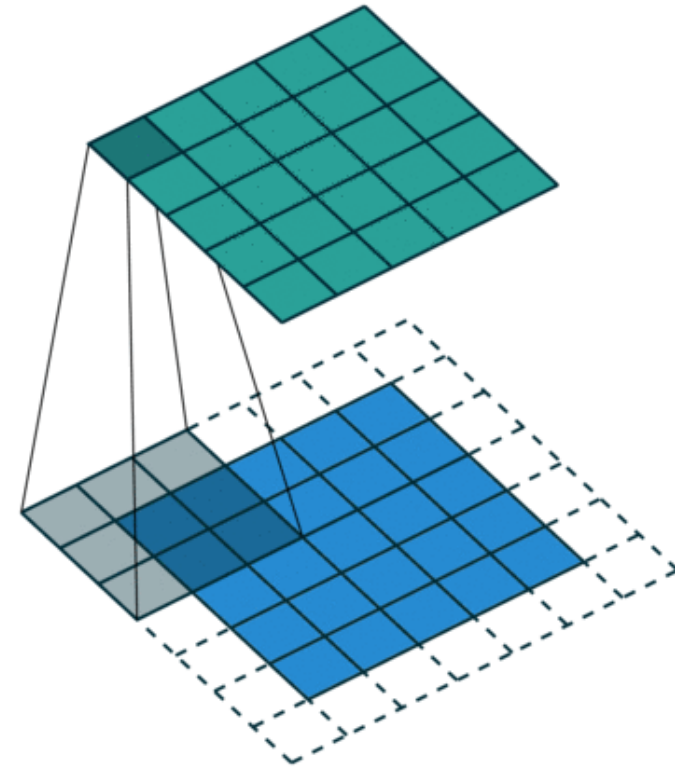


We can avoid this by padding the edges in different ways.

Padding

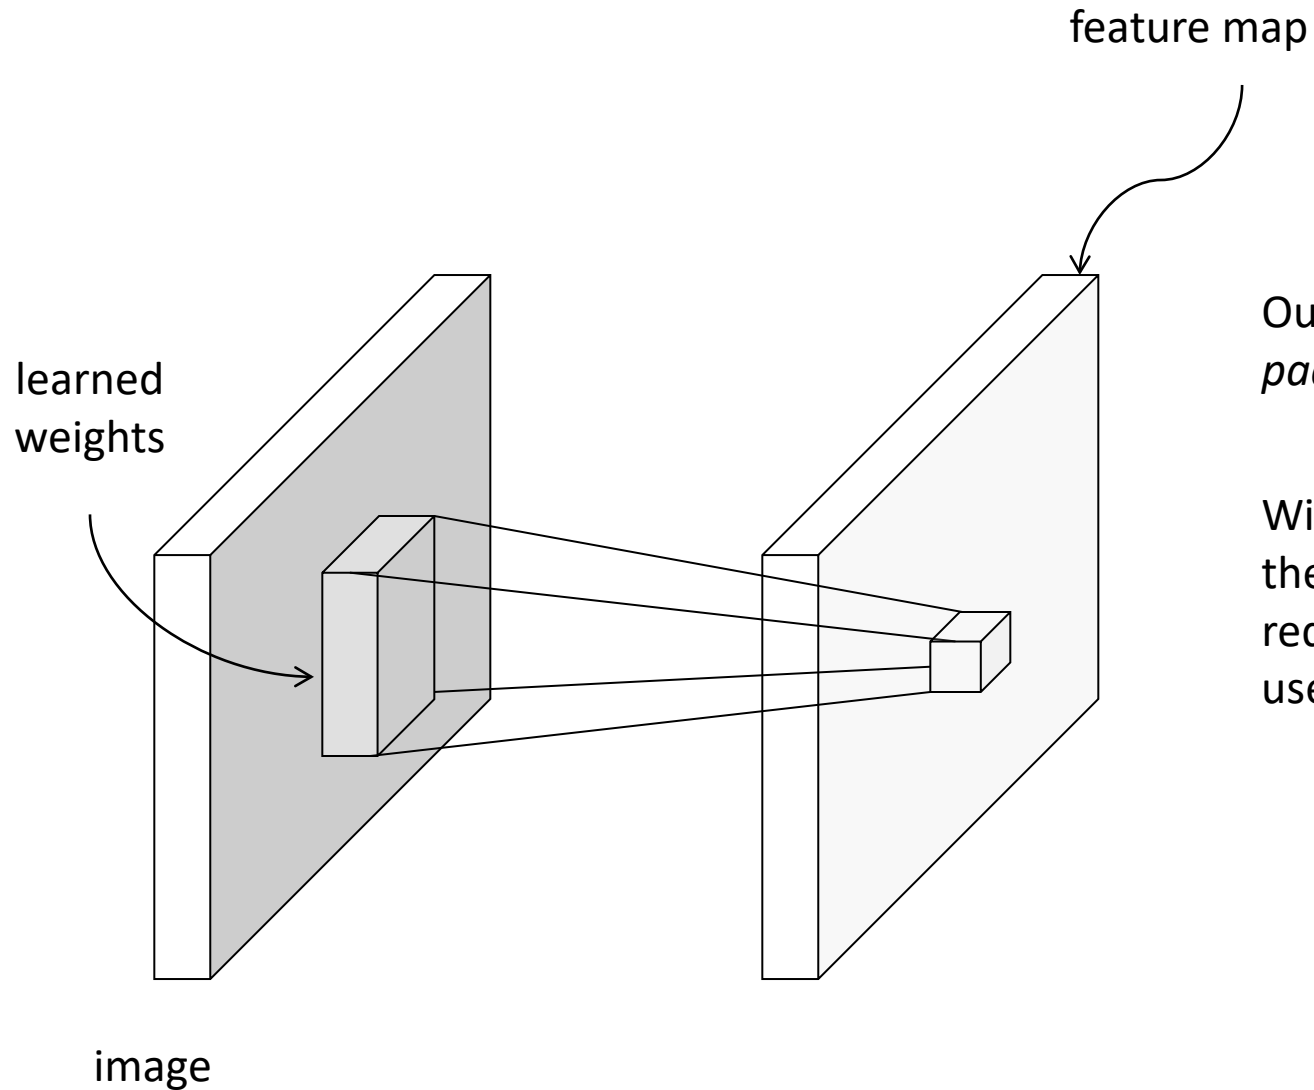


Full padding. Introduces zeros such that all pixels are visited the same amount of times by the filter. Increases size of output.



Same padding. Ensures that the output has the same size as the input.

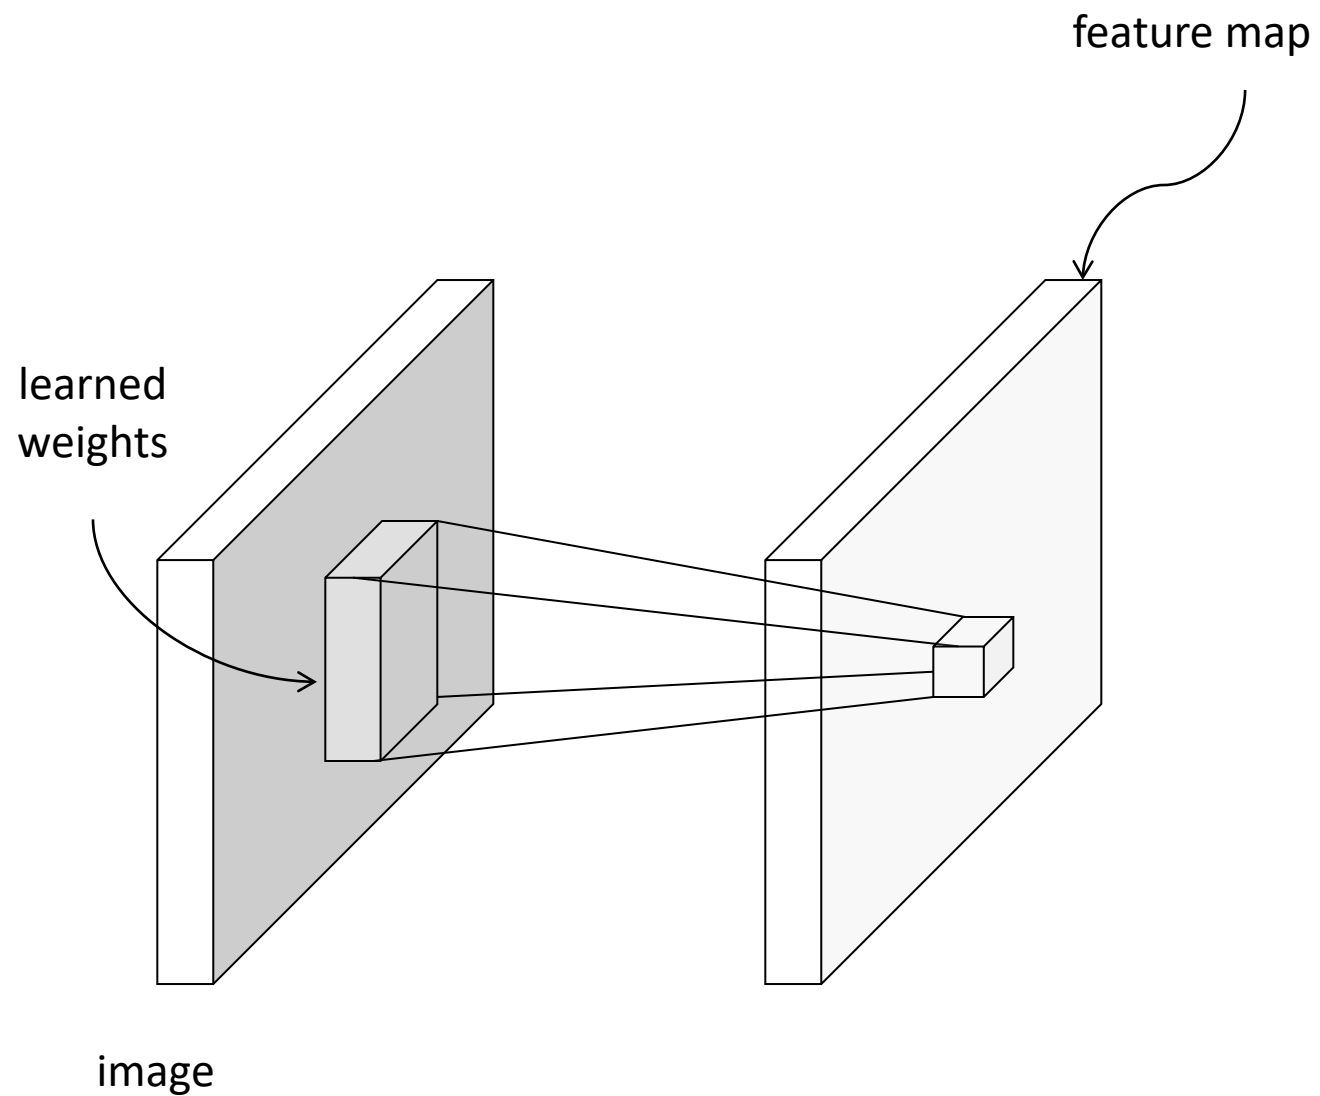
Convolutional architecture



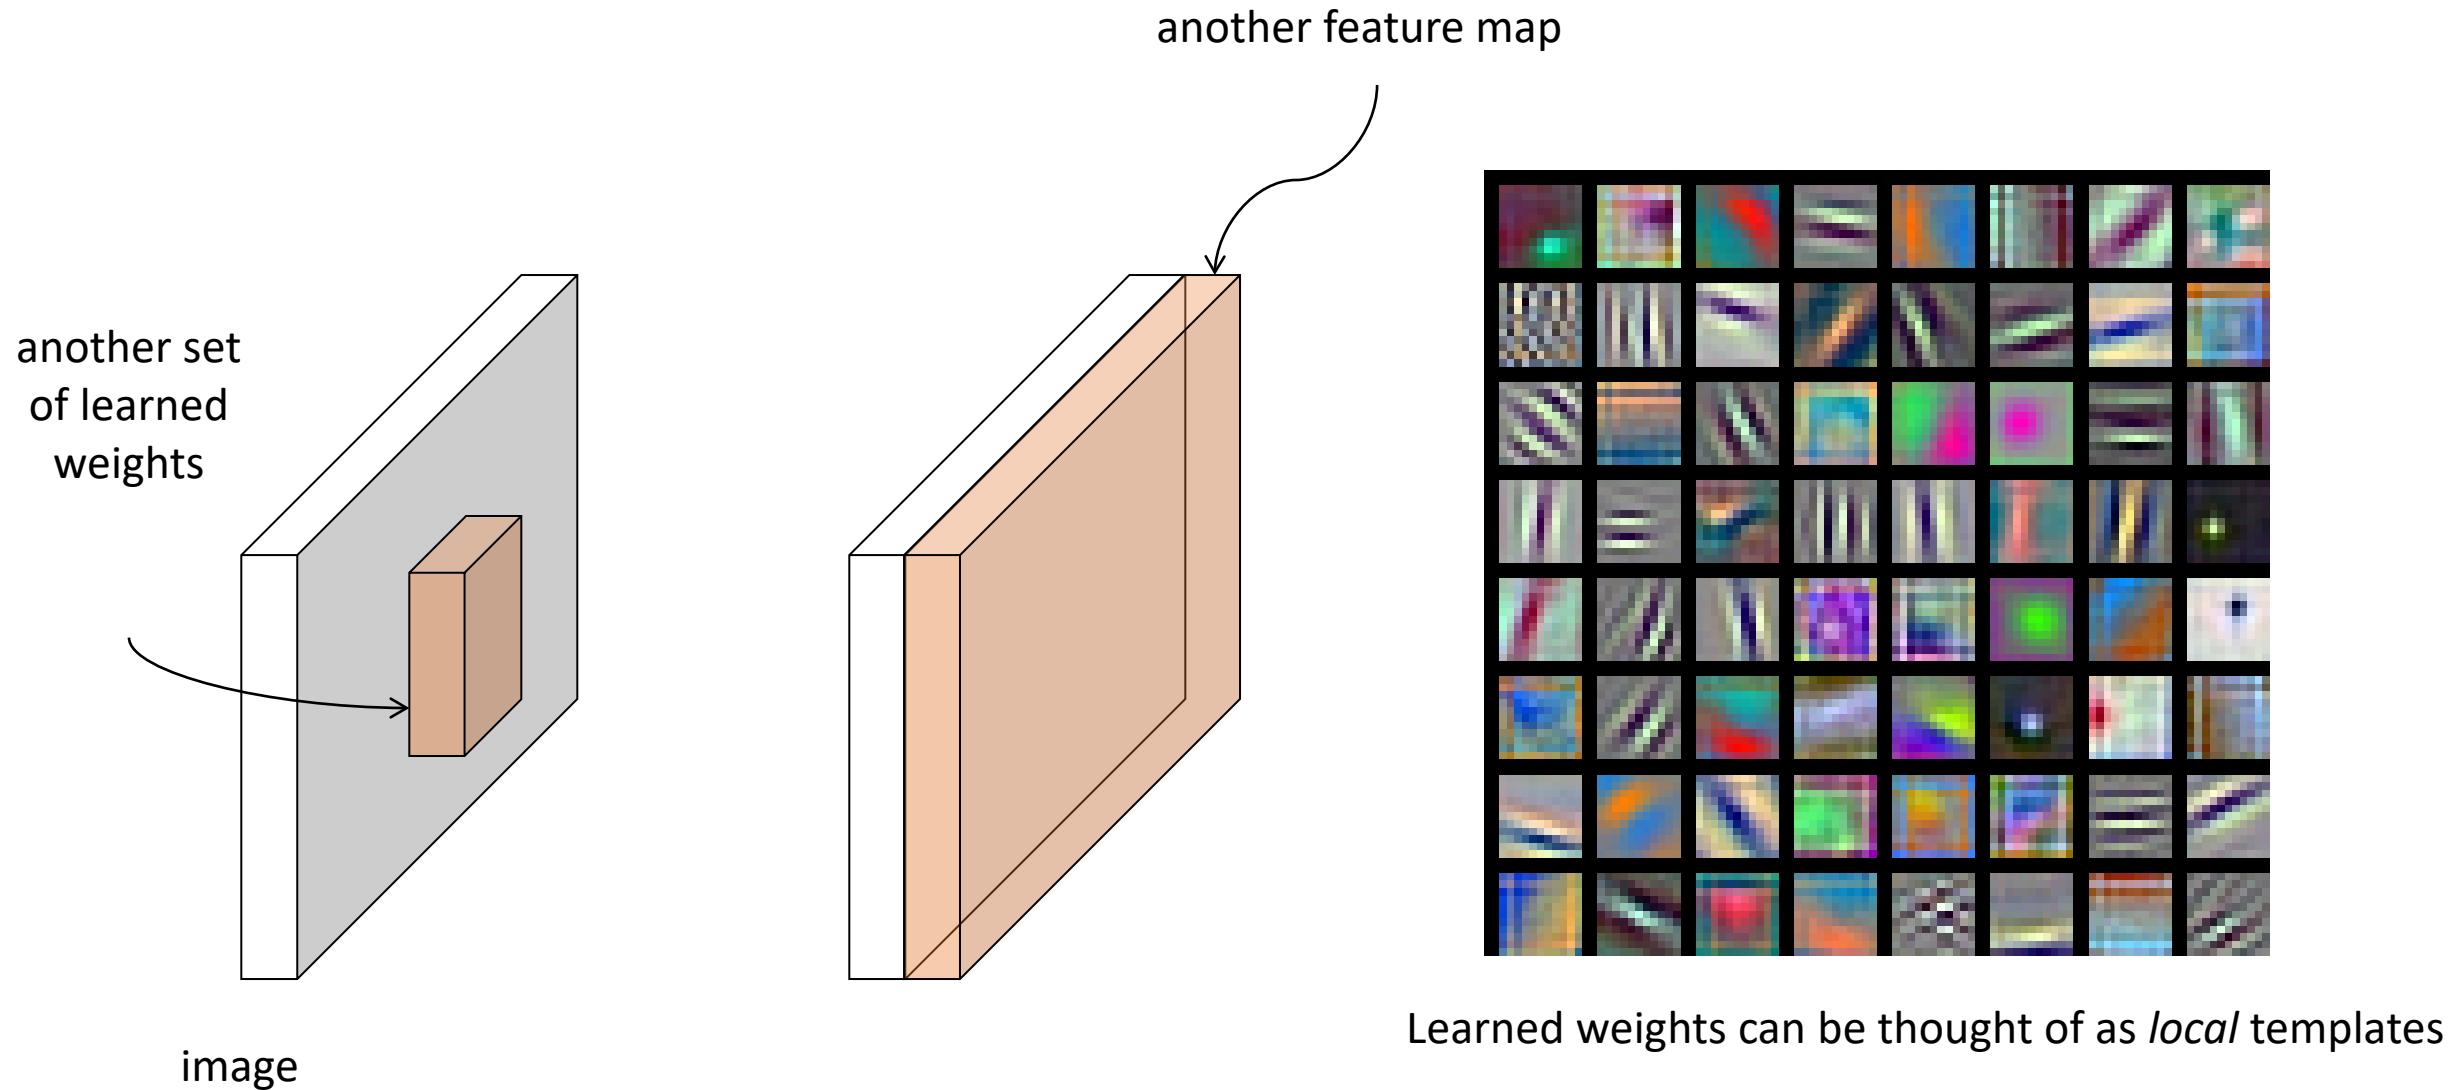
Output feature map resolution depends on *padding* and *stride*

With padding, spatial resolution remains the same if stride of 1 is used, is reduced by factor of $1/S$ if stride of S is used

Convolutional architecture

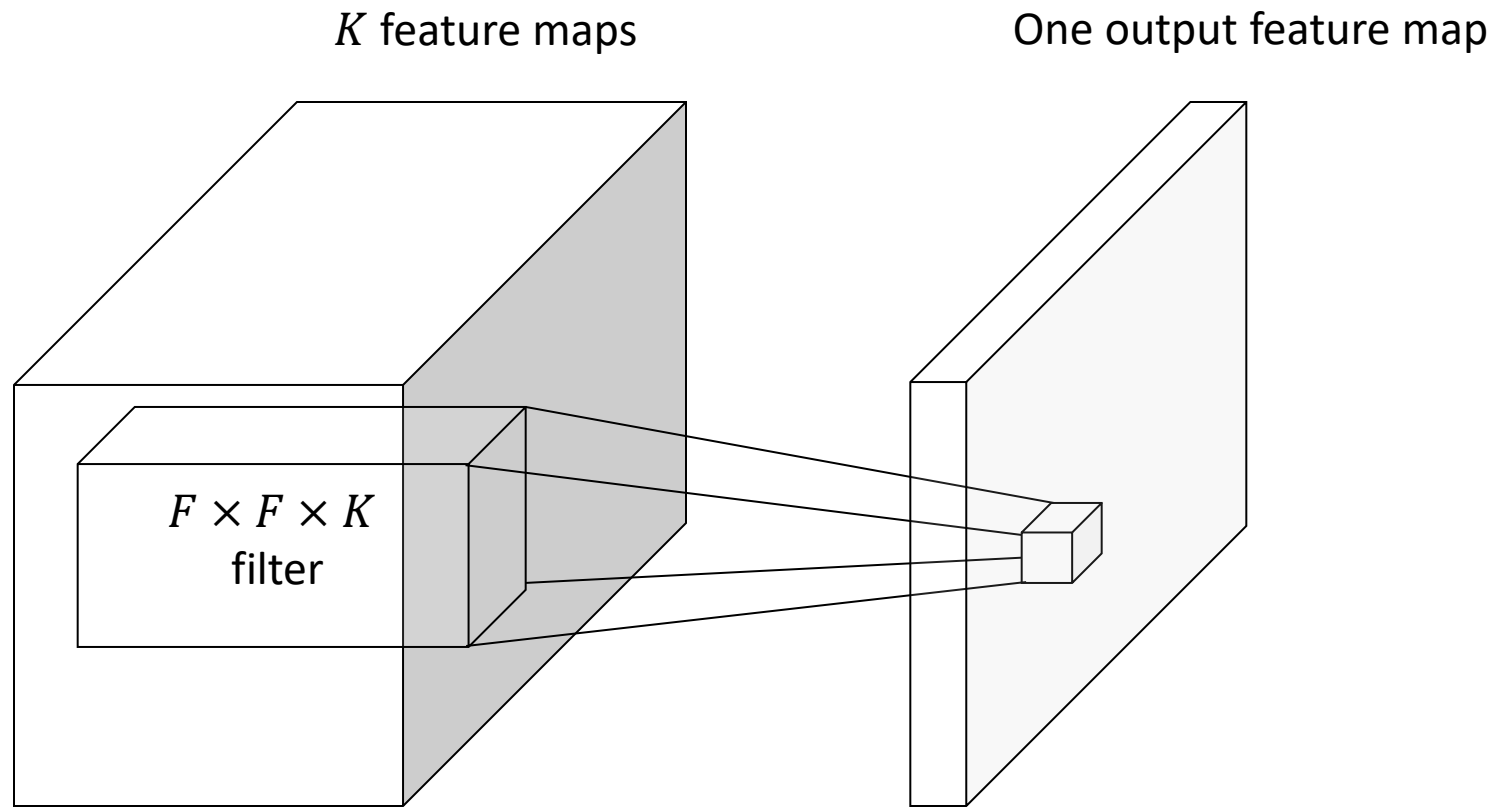


Convolutional architecture



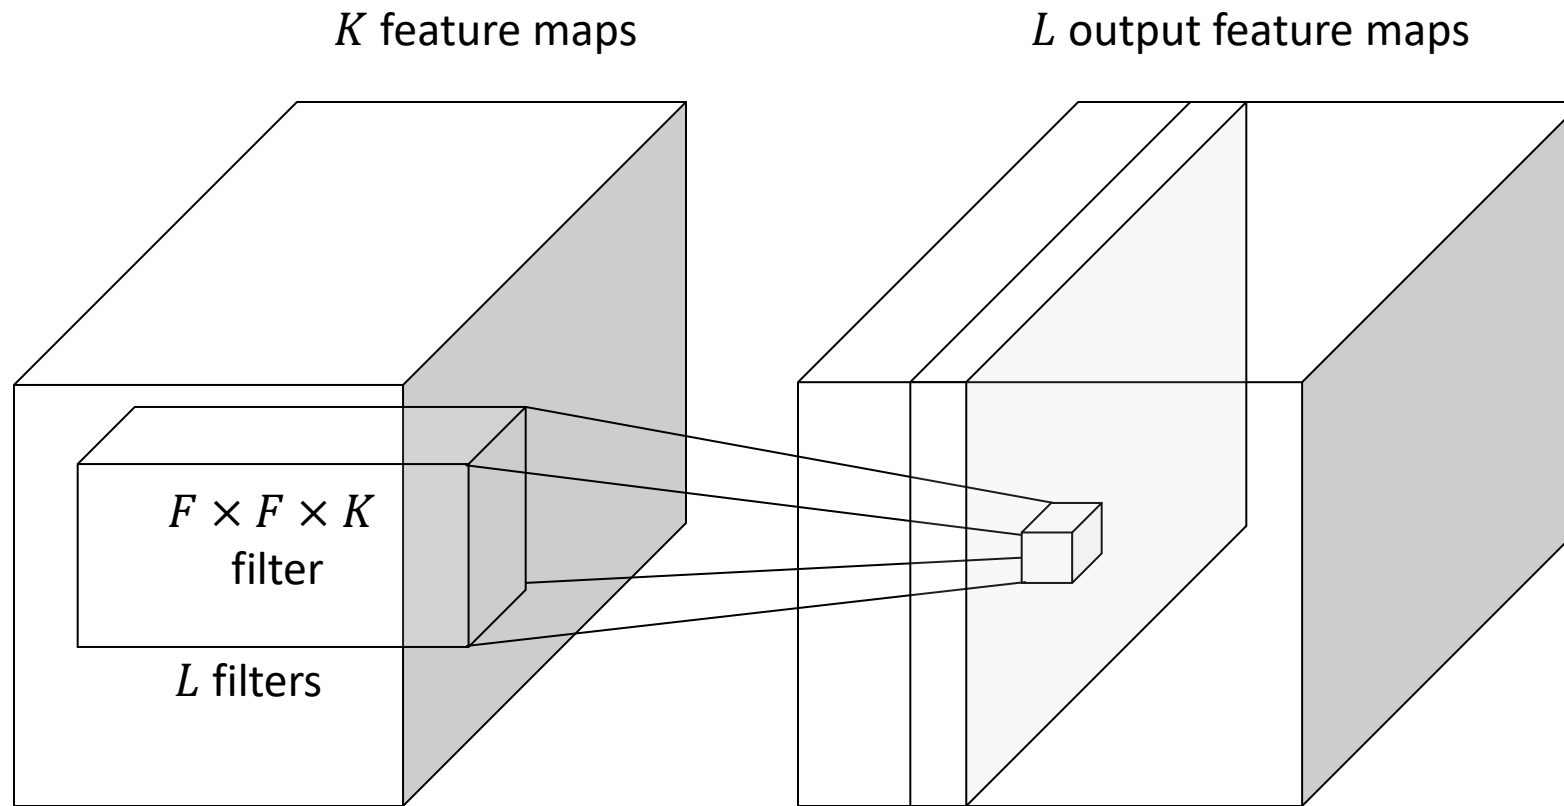
Three-dimensional convolutions

What if the *input* to a convolutional layer is a stack of K feature maps?

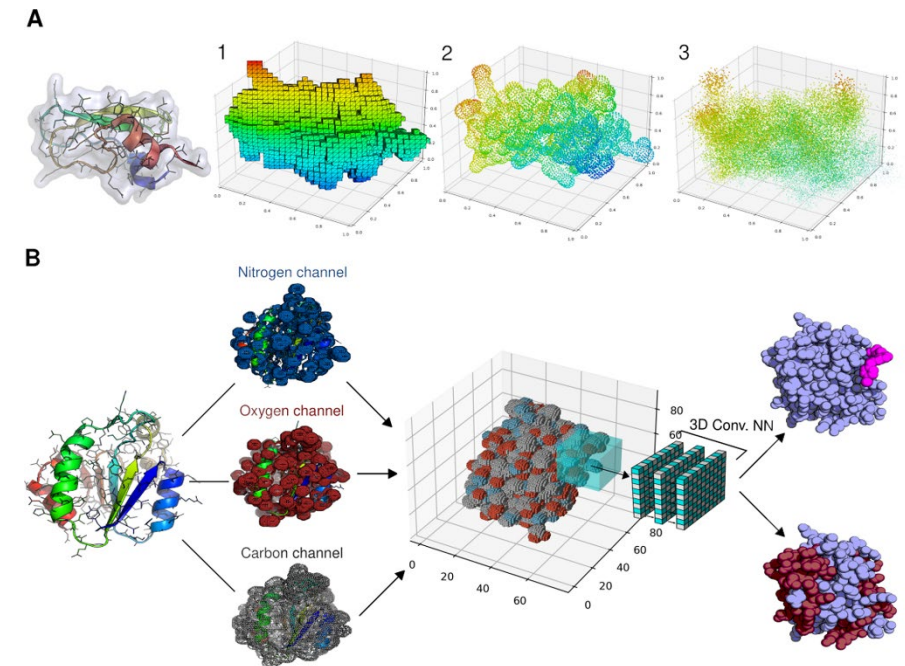
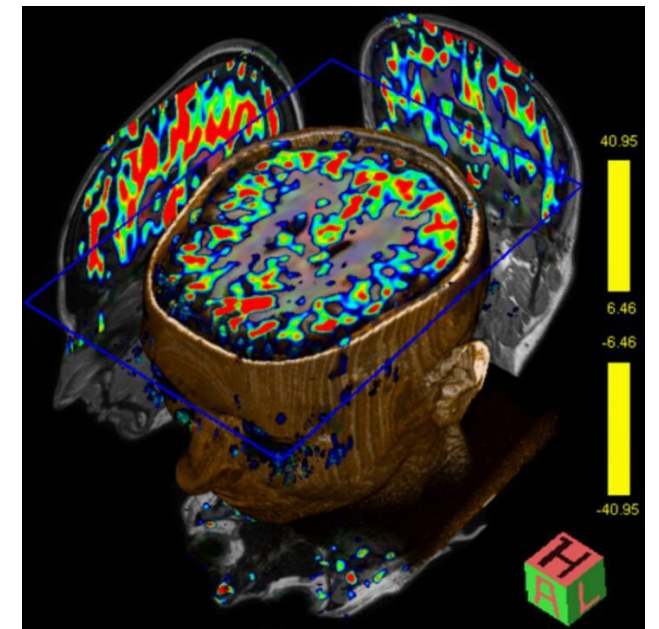
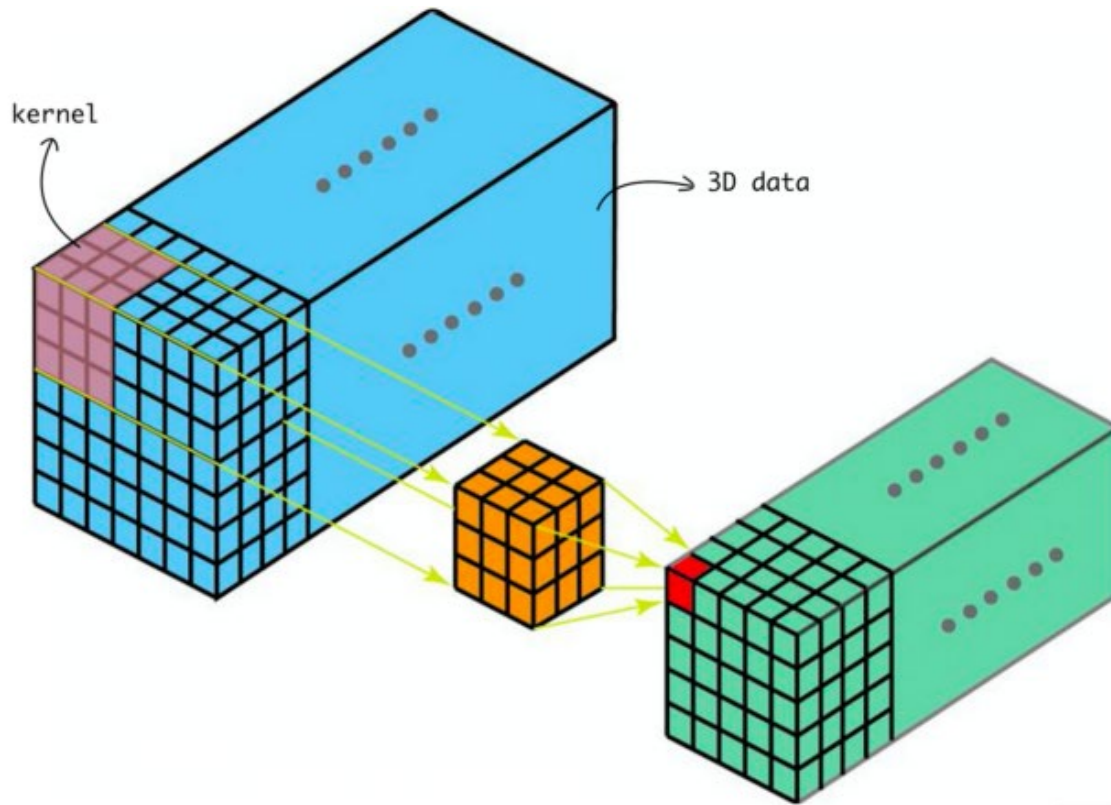


Three-dimensional convolutions

What if the *input* to a convolutional layer is a stack of K feature maps?



3-dimensional CNN aka Conv3D

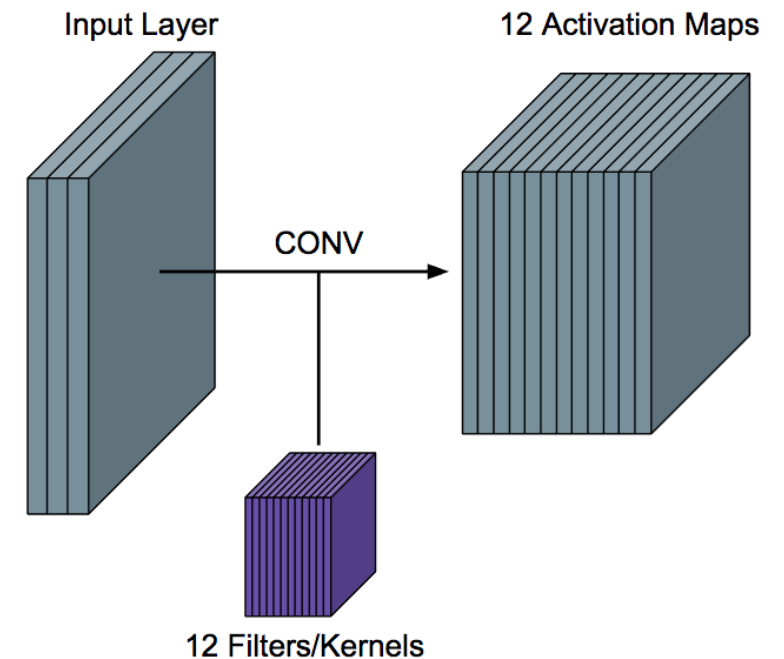


Convolutional layers (cont)

To be clear: each filter is convolved with the entirety of the **3D input cube**, but generates a **2D feature map**.

Because we have multiple filters, we end up with a 3D output: **one 2D feature map per filter**.

The feature map dimension can **change drastically** from one conv layer to the next: we can enter a layer with a 32x32x16 input and exit with a 32x32x128 output if that layer has 128 filters.



Learning CNN

In a convolutional layer, we are basically applying multiple filters at over the image to extract different features.

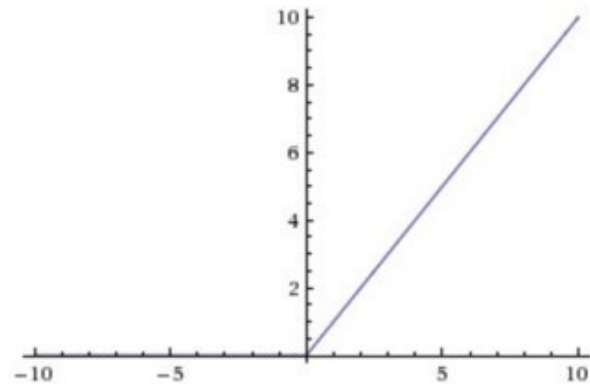
But most importantly, **we are learning those filters!**

One thing we're missing: non-linearity.

ReLU

The most successful non-linearity for CNNs is the Rectified Non-Linear unit (ReLU):

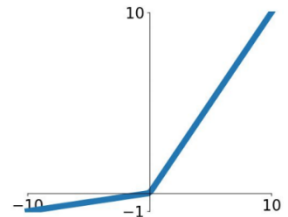
Output = $\text{Max}(\text{zero}, \text{Input})$



Combats the vanishing gradient problem occurring in sigmoids, is easier to compute, generates sparsity (not always beneficial)

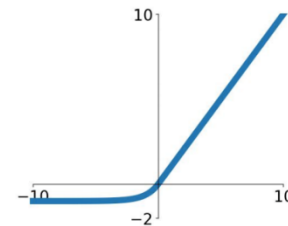
Some alternatives to ReLU:

Leaky ReLU
 $\max(0.1x, x)$

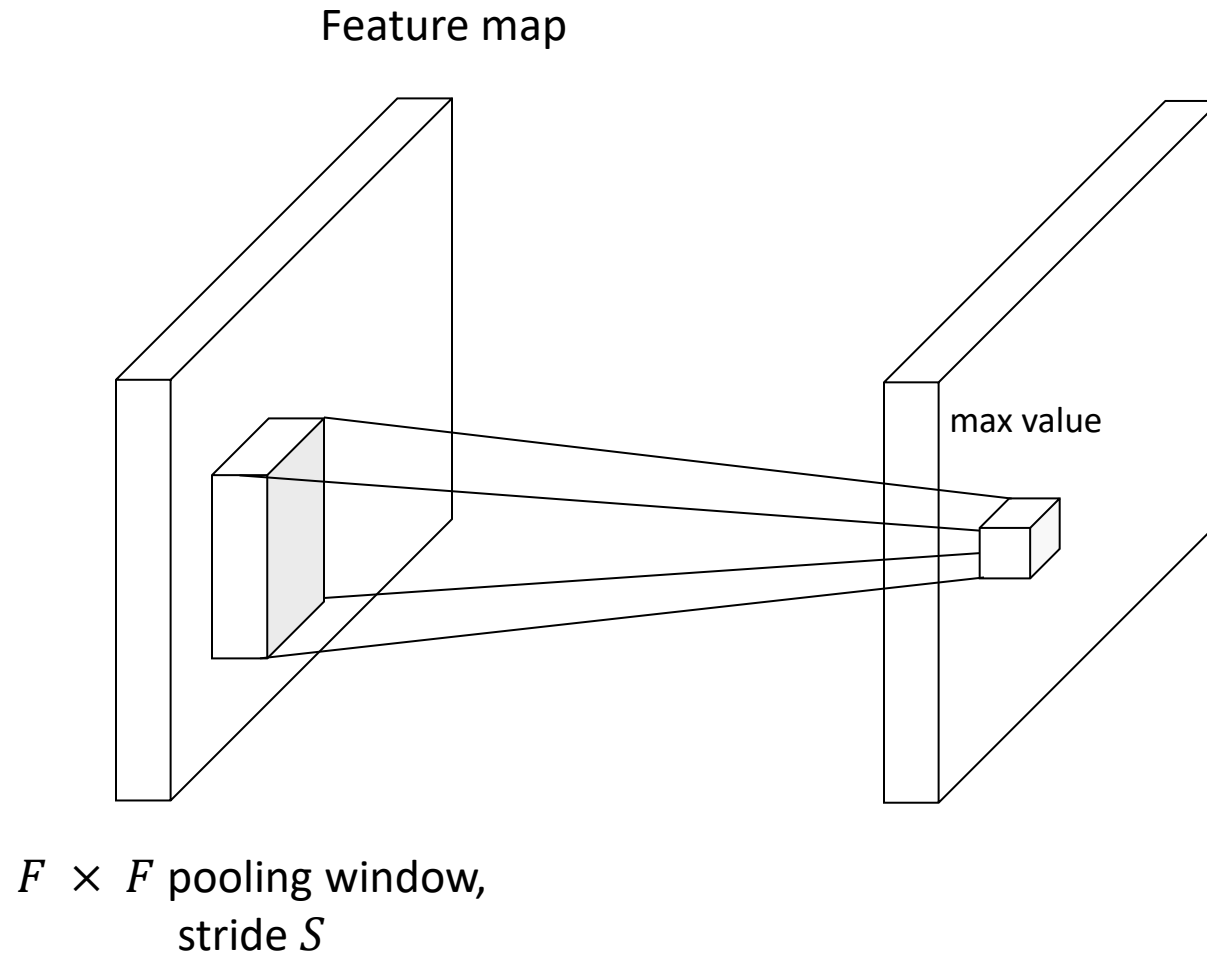


ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Max pooling layer



Usually: $F = 2$ or 3 , $S = 2$

Max pooling: Example

Single channel

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with 2×2
kernel size and stride 2



Max pooling: Example

Single channel

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with 2×2
kernel size and stride 2



Max pooling: Example

Single channel

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Max pooling with 2×2
kernel size and stride 2



6	8
3	4

Convolutional layers so far

- A convolutional layer convolves each of its filters with the input.
- Input: a **3D tensor**, where the dimensions are Width, Height and Channels (or Feature Maps)
- Output: a **3D tensor**, with dimensions Width, Height and Feature Maps (one for each filter)
- Applies non-linear **activation function** (usually ReLU) over each value of the output.
- Multiple **parameters to define**: number of filters, size of filters, stride, padding, activation function to use, regularization.

Building a CNN

A convolutional neural network is built by stacking layers, typically of 3 types:



Convolutional
Layers

Pooling Layers

Fully connected
Layers

Building a CNN

- A
3

Convolutional Layers

Action

- Apply filters to extract features
- Filters are composed of small kernels, learned.
- One bias per filter.
- Apply activation function on every value of feature map

Parameters

- Number of kernels
- Size of kernels (W and H only, D is defined by input cube)
- Activation function
- Stride
- Padding
- Regularization type and value

I/O

- Input: 3D cube, previous set of feature maps
- Output: 3D cube, one 2D map per filter

Building a CNN

- A convolutional neural network is built by stacking layers, typically of 3 types:

Convolutional
Layers

Pooling Layers

Fully connected
Layers

Building a CNN

- A 3

Pooling Layers

Action

- Reduce dimensionality
- Extract maximum of average of a region
- Sliding window approach

Parameters

- Stride
- Size of window

I/O

- Input: 3D cube, previous set of feature maps
- Output: 3D cube, one 2D map per filter, reduced spatial dimensions



Building a CNN

- A convolutional neural network is built by stacking layers, typically of 3 types:

Convolutional
Layers

Pooling Layers

Fully connected
Layers

Building a CNN

- A 3

Fully connected Layers

Action

- Aggregate information from final feature maps
- Generate final classification

Parameters

- Number of nodes
- Activation function: usually changes depending on role of layer. If aggregating info, use ReLU. If producing final classification, use Softmax.

I/O

- Input: FLATTENED 3D cube, previous set of feature maps
- Output: 3D cube, one 2D map per filter

of

Examples

- I have a convolutional layer with 16 3x3 filters that takes an RGB image as input.
 - What else can we define about this layer?
 - Activation function
 - Stride
 - Padding type
 - How many parameters does the layer have?

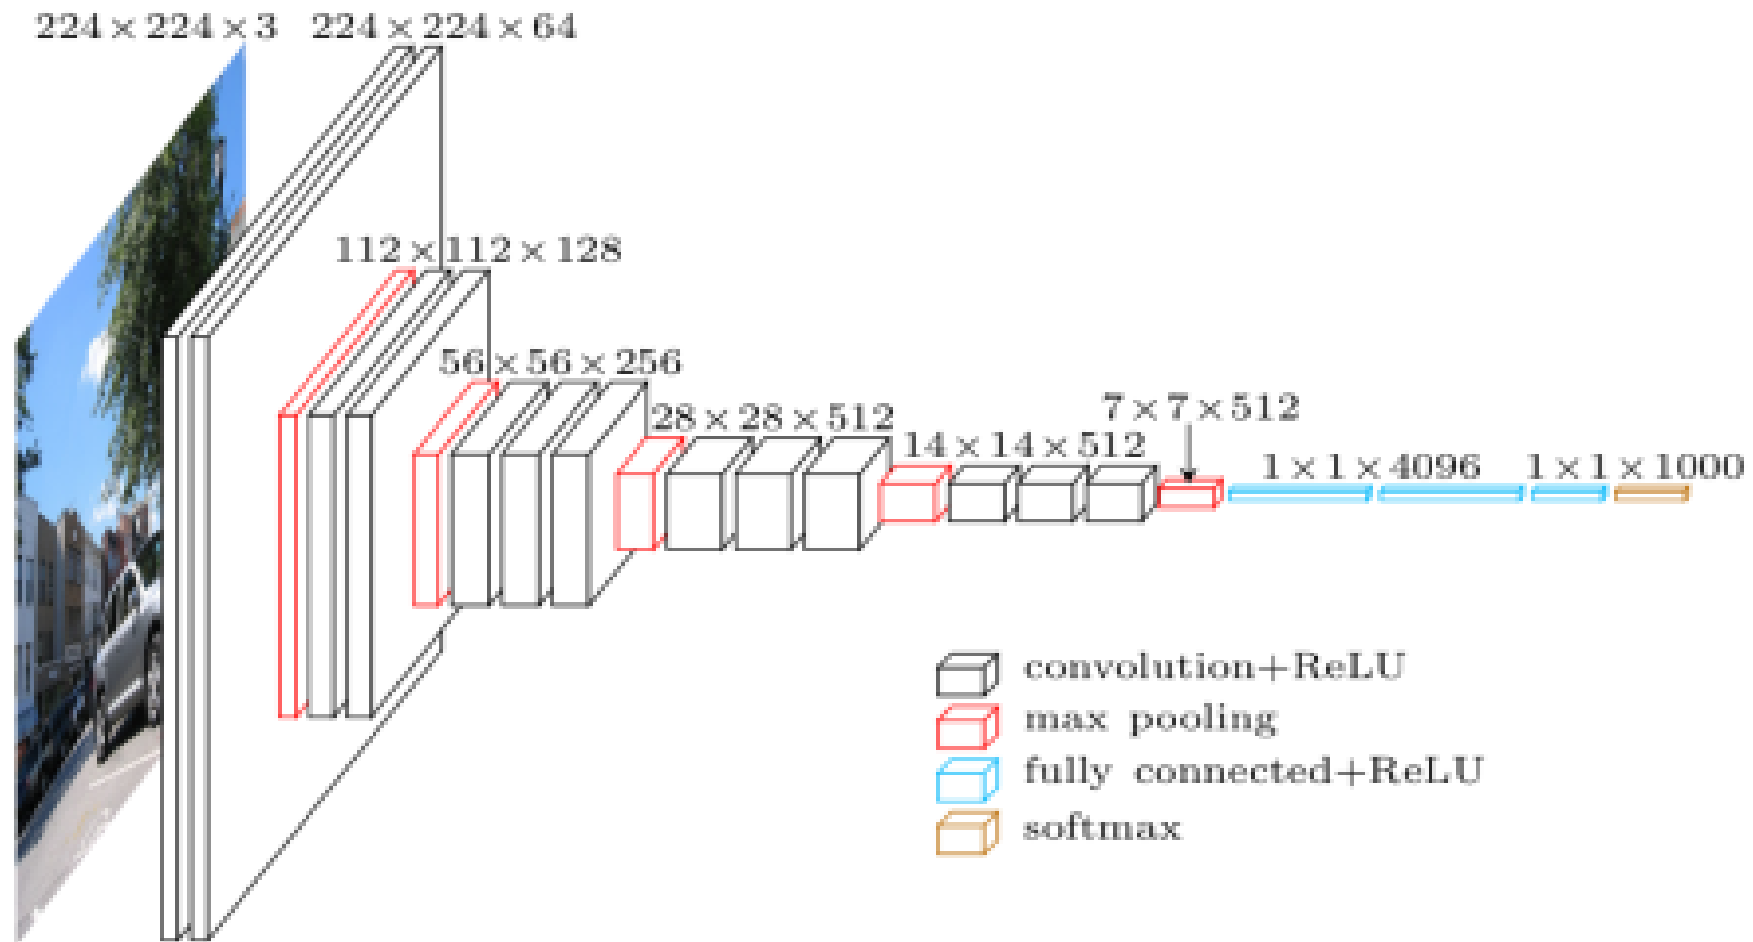
$$16 \times 3 \times 3 \times 3 + 16 = 448$$

Number of
filters

Size of
Filters
Number of
channels of
prev layer

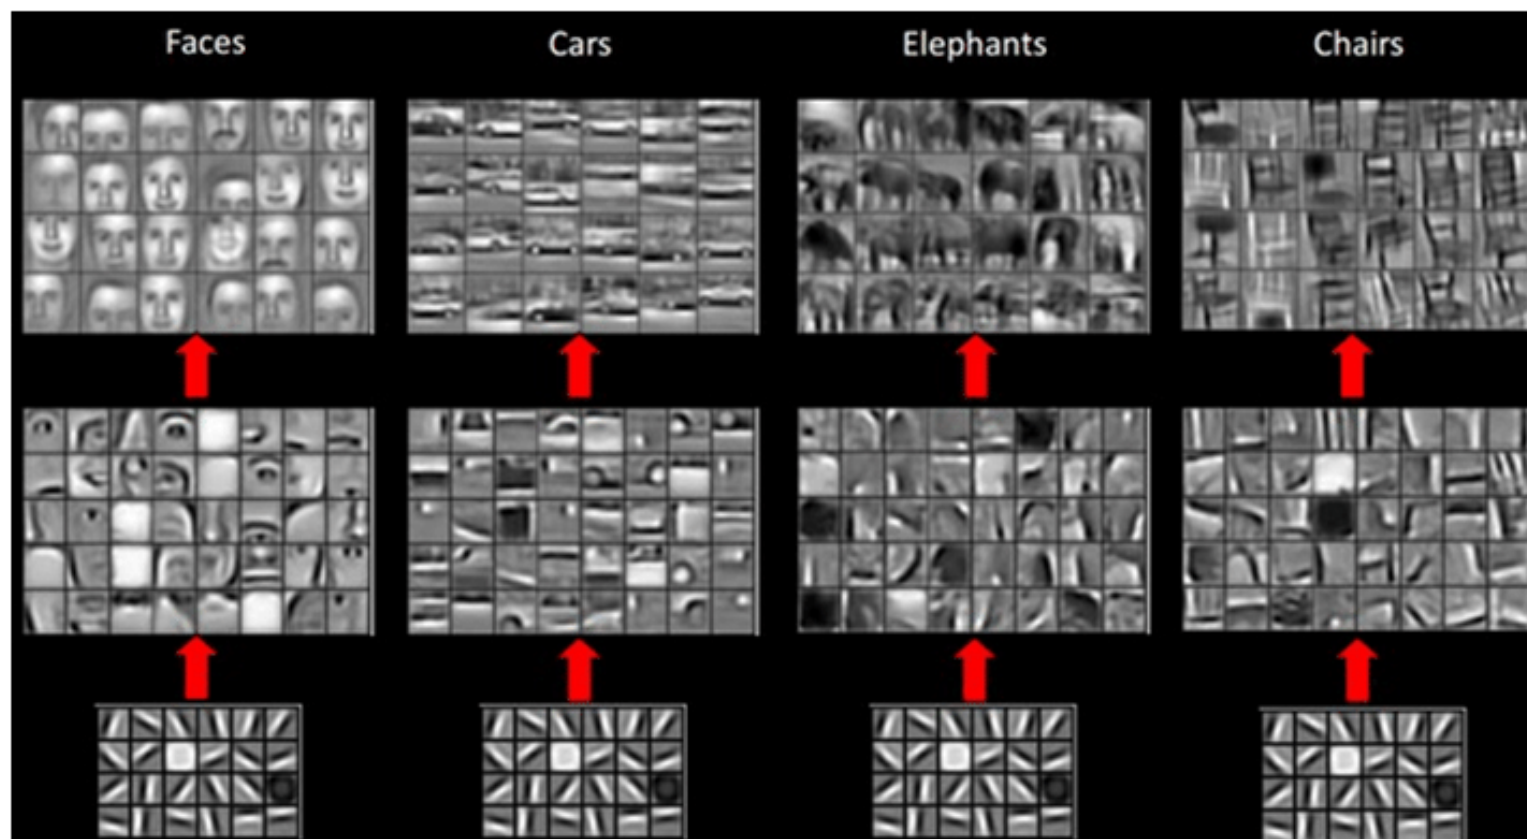
Biases (one
per filter)

Fully built CNN (VGG)



What do CNN layers learn?

- Each CNN layer learns filters of increasing complexity.
- The first layers learn **basic feature detection filters**: edges, corners, etc.
- The middle layers learn filters that detect **parts of objects**. For faces, they might learn to respond to eyes, noses, etc.
- The last layers have higher representations: they learn to **recognize full objects**, in different shapes and positions.

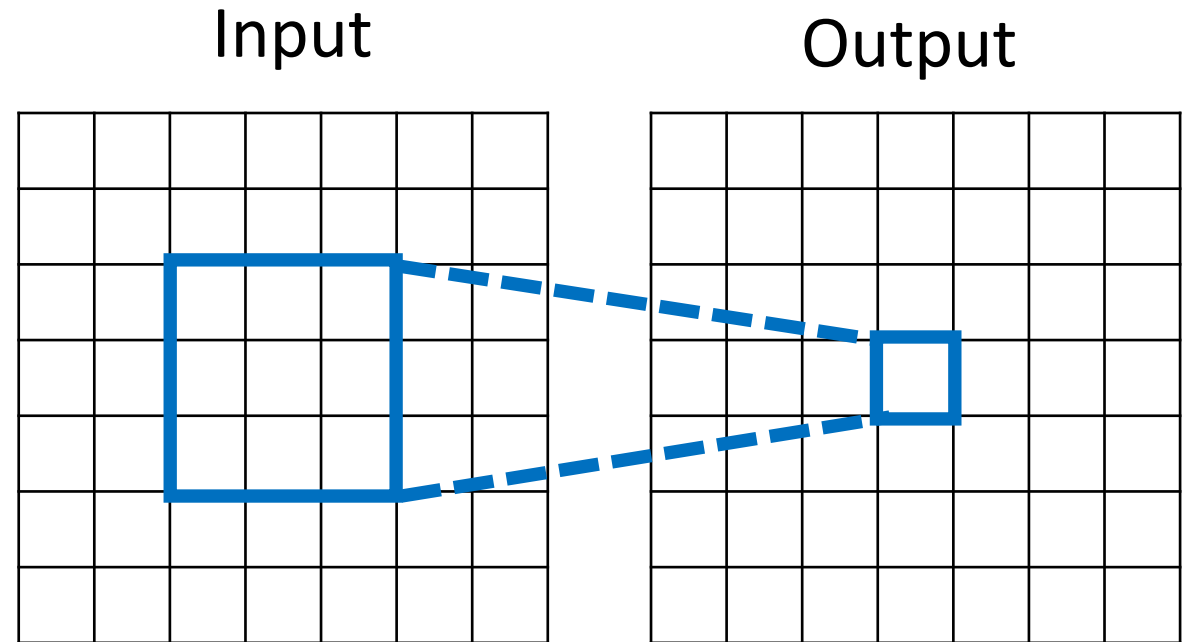


- 3D visualization of networks in action
- <http://scs.ryerson.ca/~aharley/vis/conv/>
- <https://www.youtube.com/watch?v=3JQ3hYko51Y>

Receptive field

3x3 convolutions, stride 1

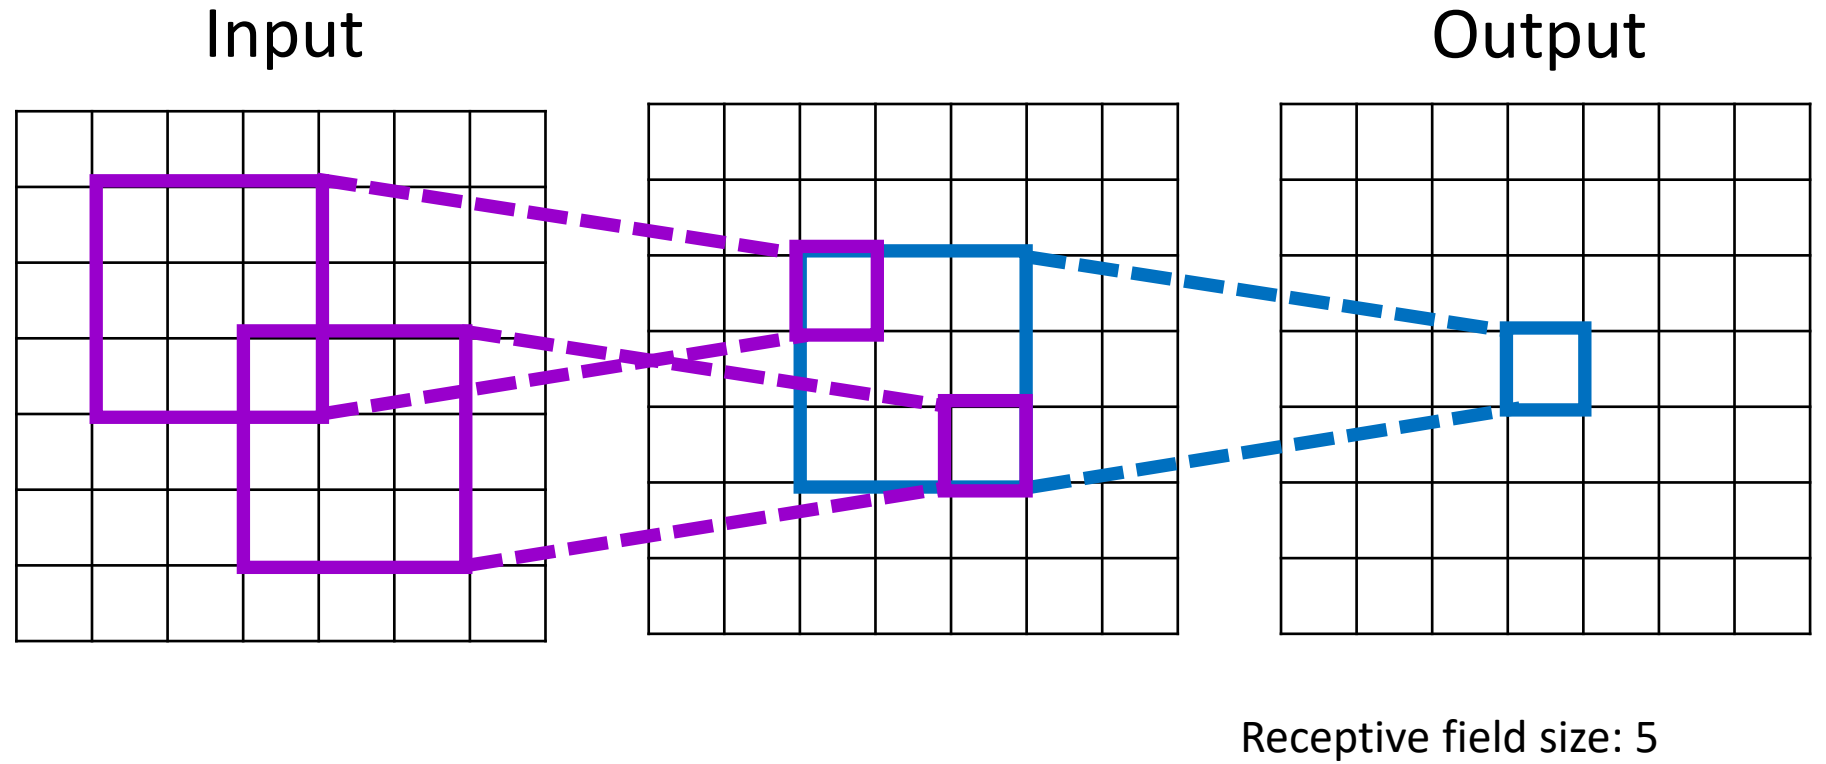
The *receptive field* of a unit is the region of the input feature map whose values contribute to the response of that unit (either in the previous layer or in the initial image)



Receptive field size: 3

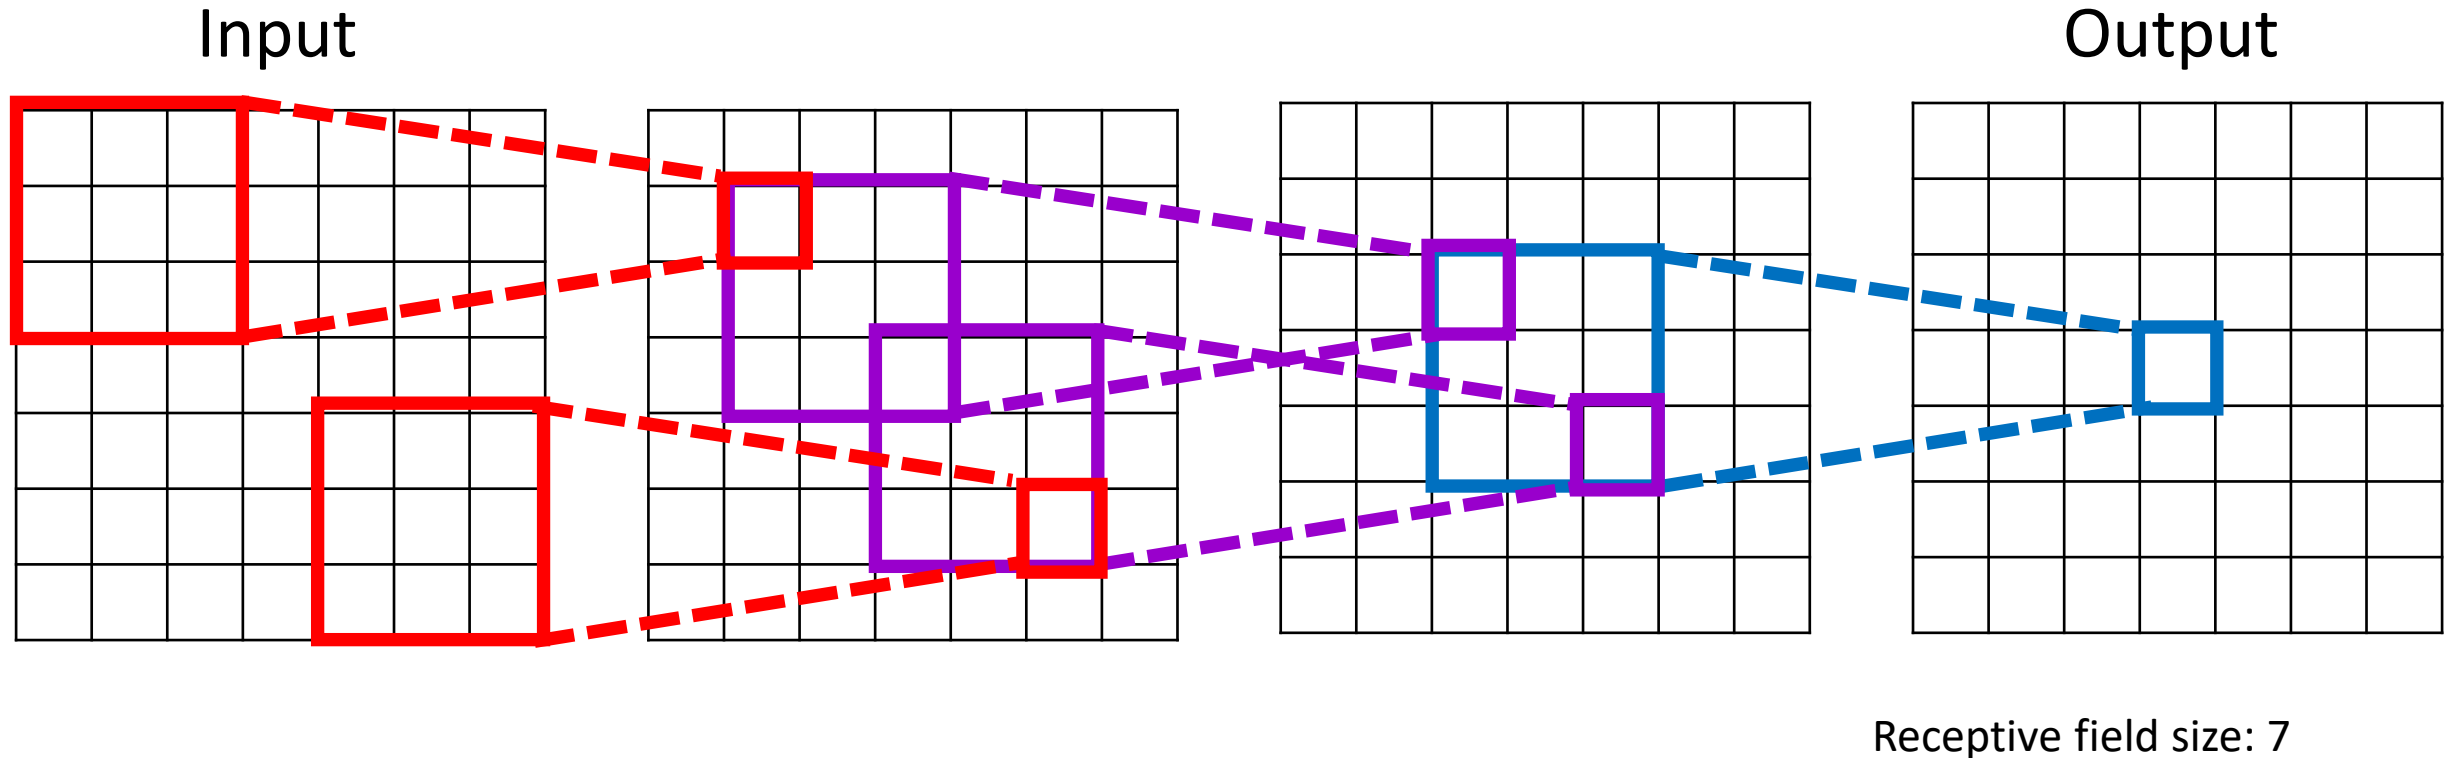
Receptive field

3x3 convolutions, stride 1



Receptive field

3x3 convolutions, stride 1

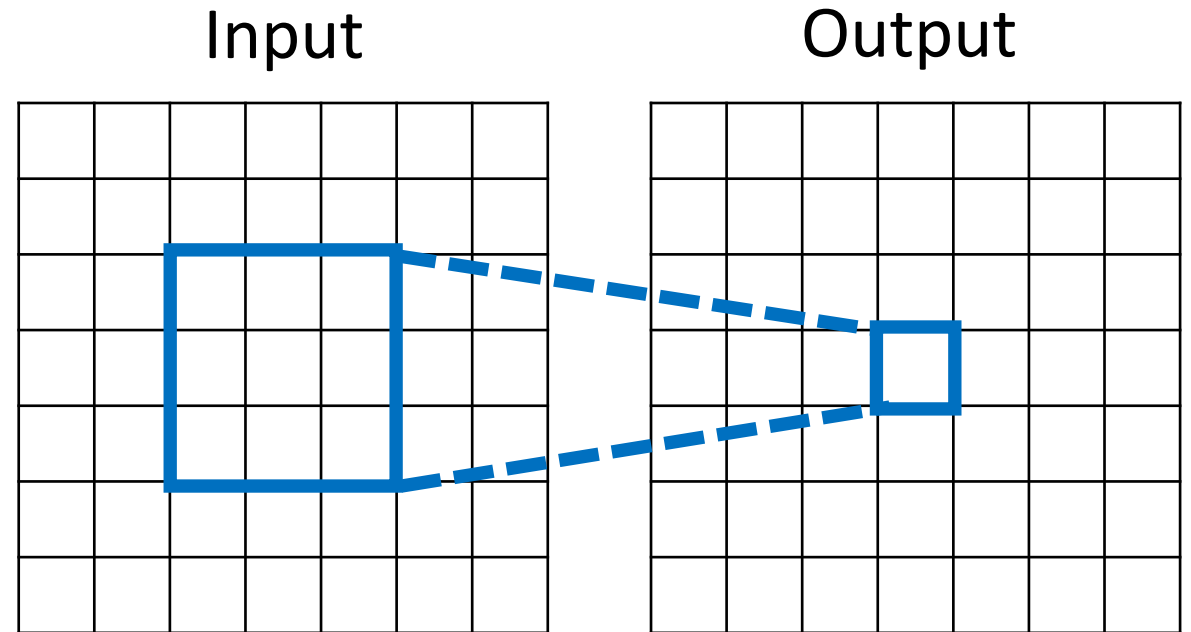


Each successive convolution adds $F - 1$ to the receptive field size

With L layers the receptive field size is $1 + L * (F - 1)$

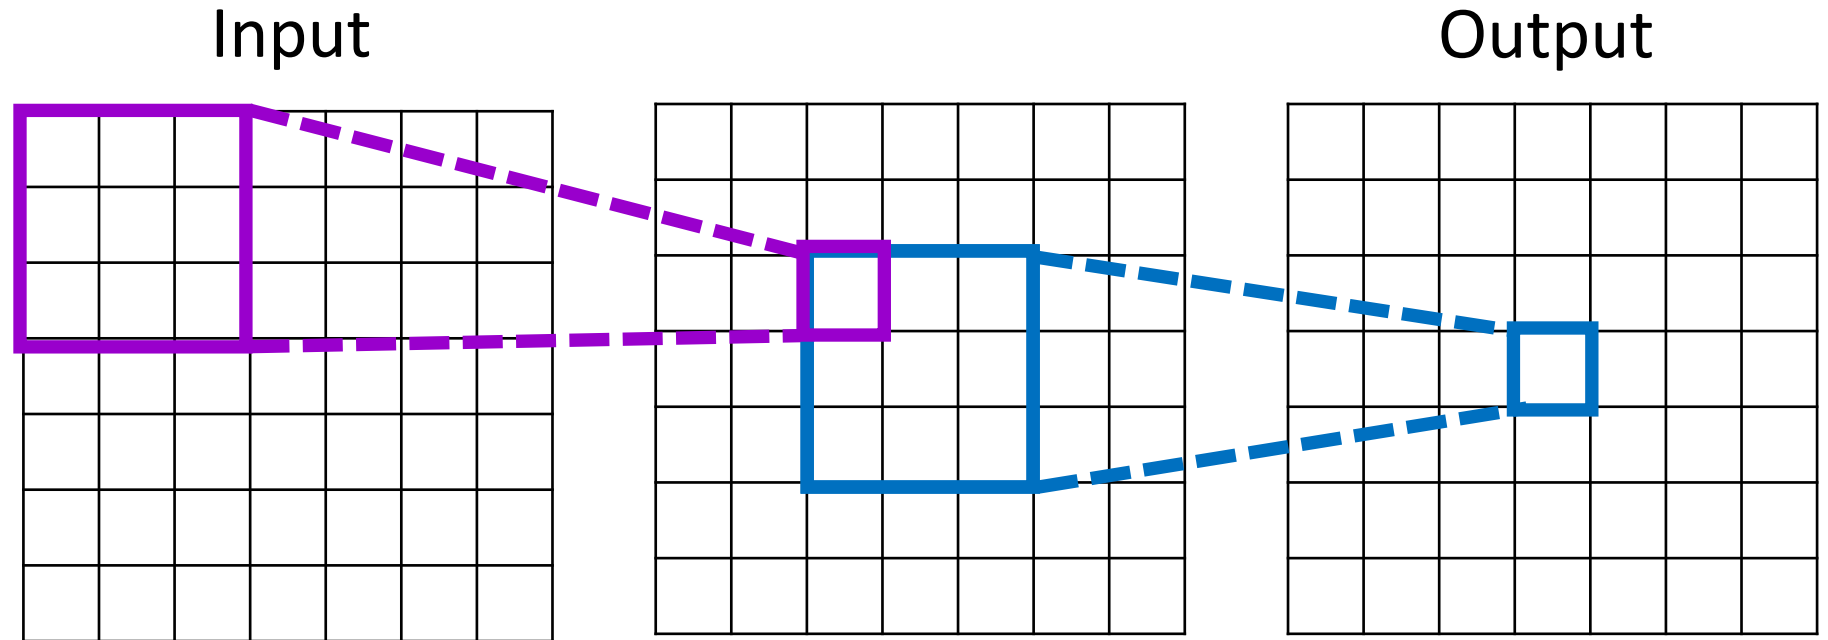
Receptive field

3x3 convolutions, stride 2



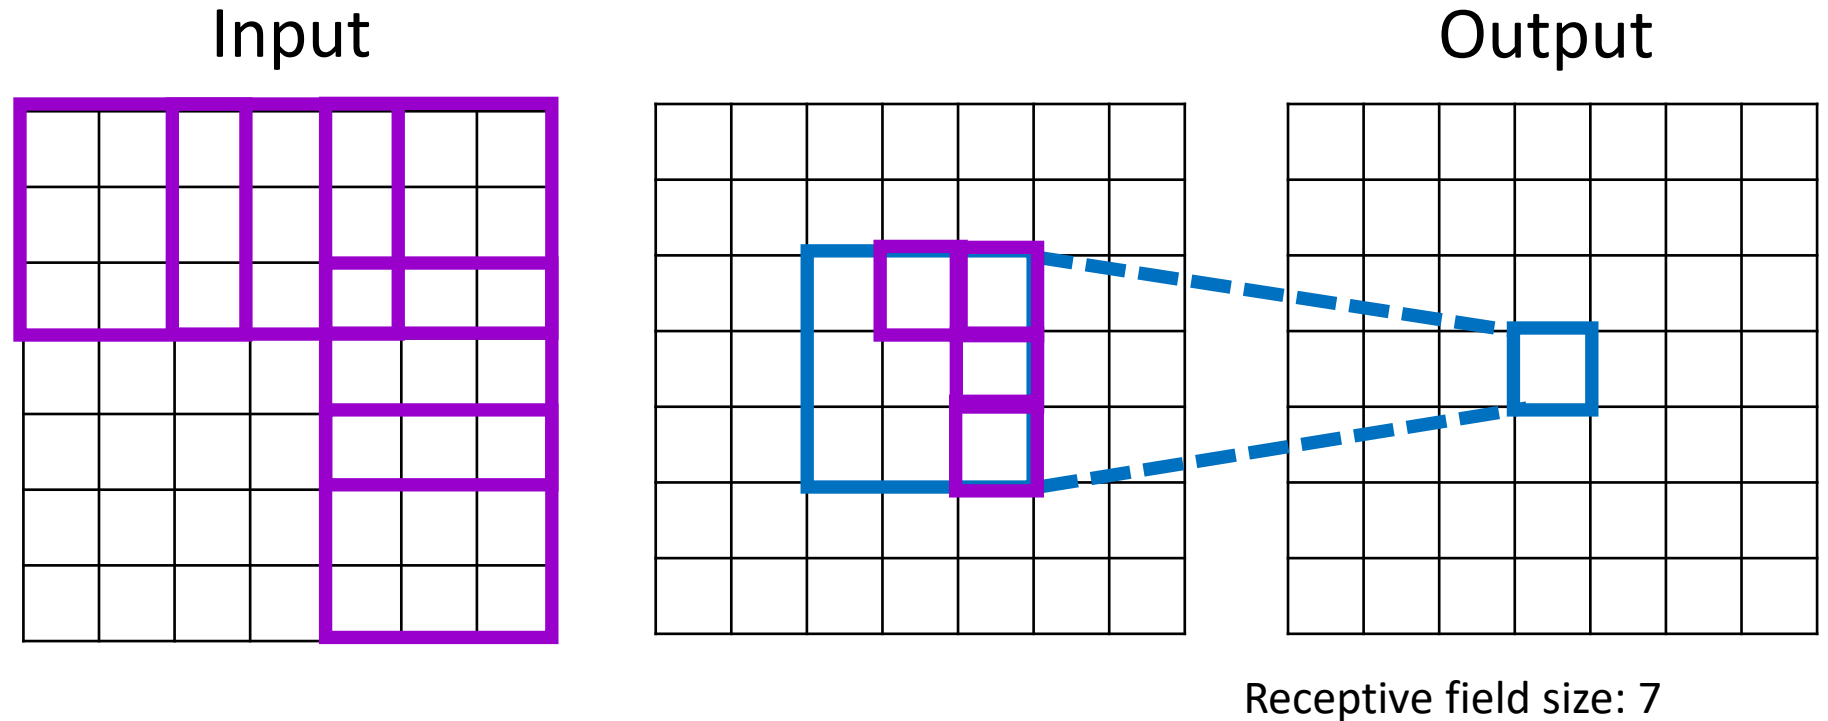
Receptive field

3x3 convolutions, stride 2



Receptive field

3x3 convolutions, stride 2



With a stride of 2, receptive field size is given by $2^{L+1} - 1$, i.e., it grows exponentially (though spatial resolution decreases exponentially)

Dropout, Overfitting & Normalization

Overfitting in Deep Neural Nets

Deep nets have many non-linear hidden layers.

Making them very expressive to learn complicated relationships between inputs and outputs

But with limited training data, many complicated relationships will be the result of training noise

So they will exist in the training set and not in test set even if drawn from same distribution

Many methods developed to reduce overfitting

Early stopping with a validation set

Weight penalties (L1 and L2 regularization)

Batch Normalization

- Training time:
 - Mini-batch of activations for layer to normalize

$$H' = \frac{H - \mu}{\sigma}$$

where

$$\mu = \frac{1}{m} \sum_i H_{i,:}$$

Vector of mean activations
across mini-batch

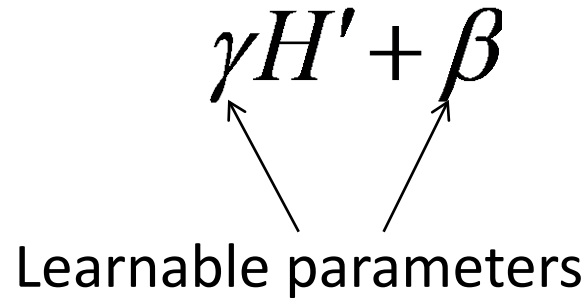
$$\sigma = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Vector of SD of each unit
across mini-batch

Batch Normalization

Training time:

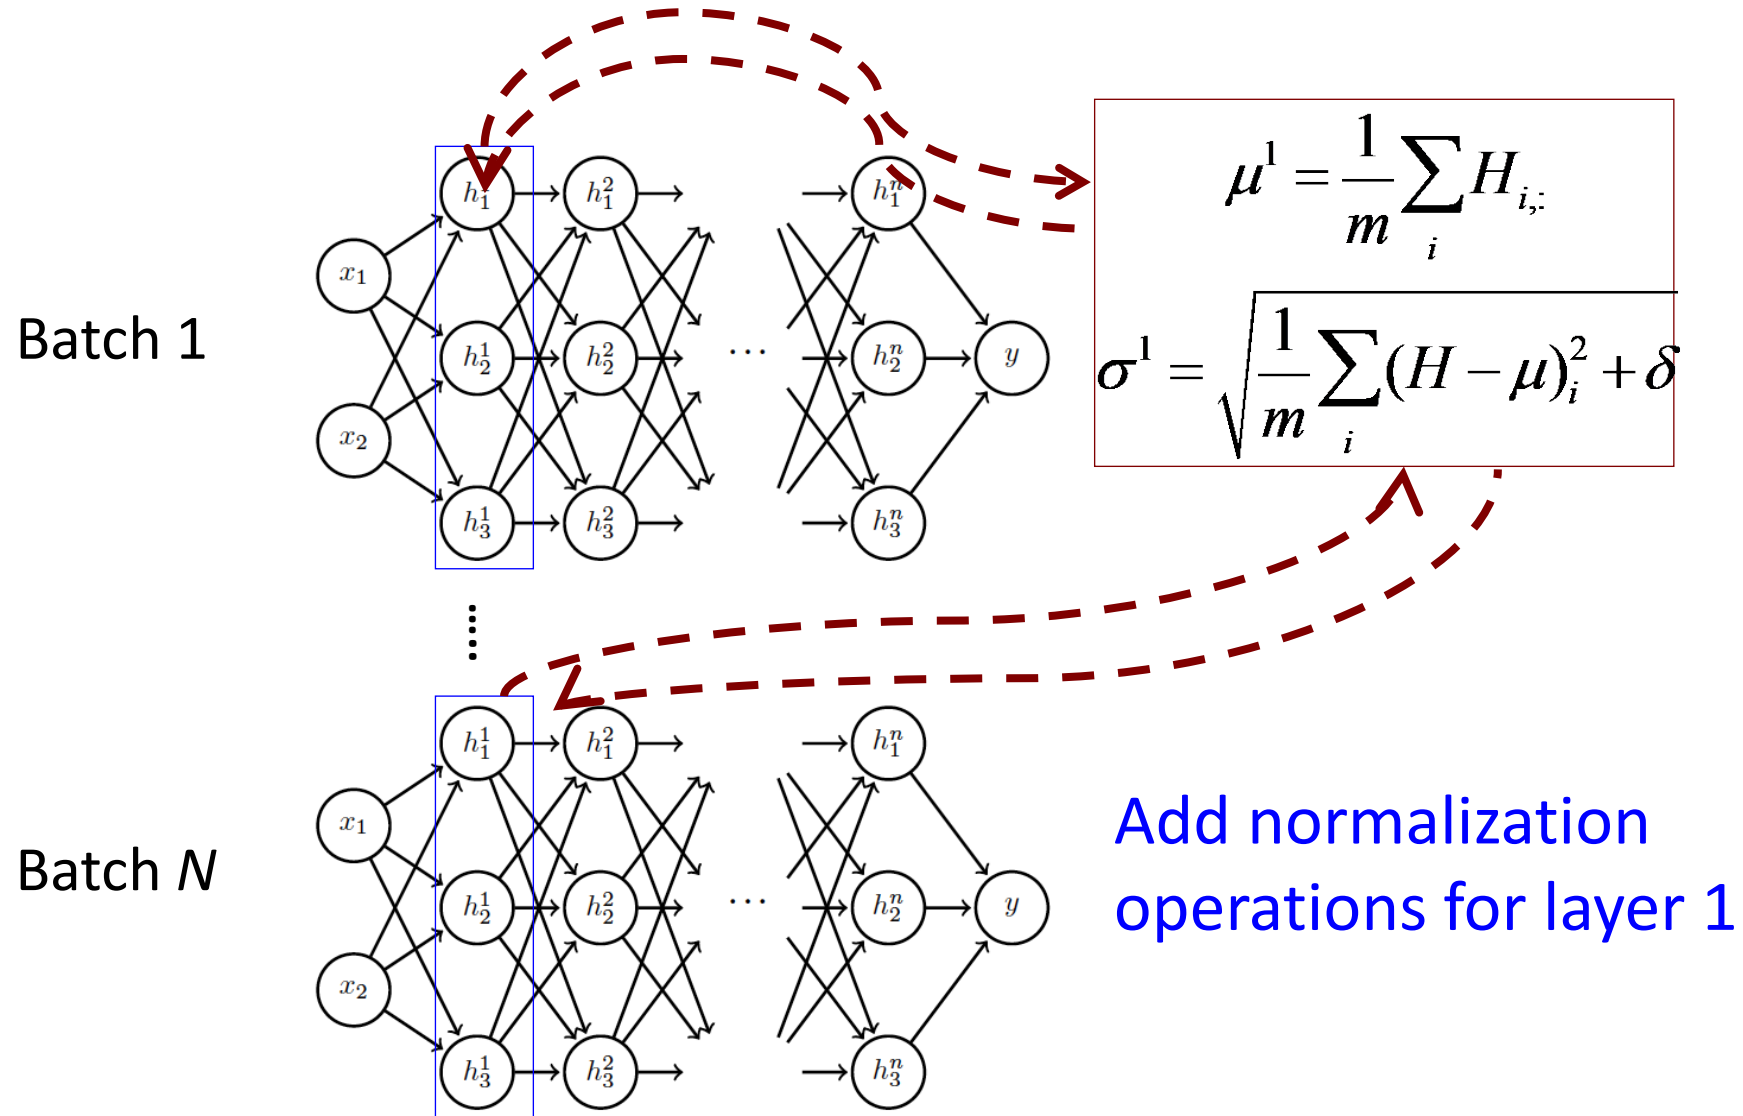
- Normalization can reduce expressive power
- Instead use:

$$\gamma H' + \beta$$


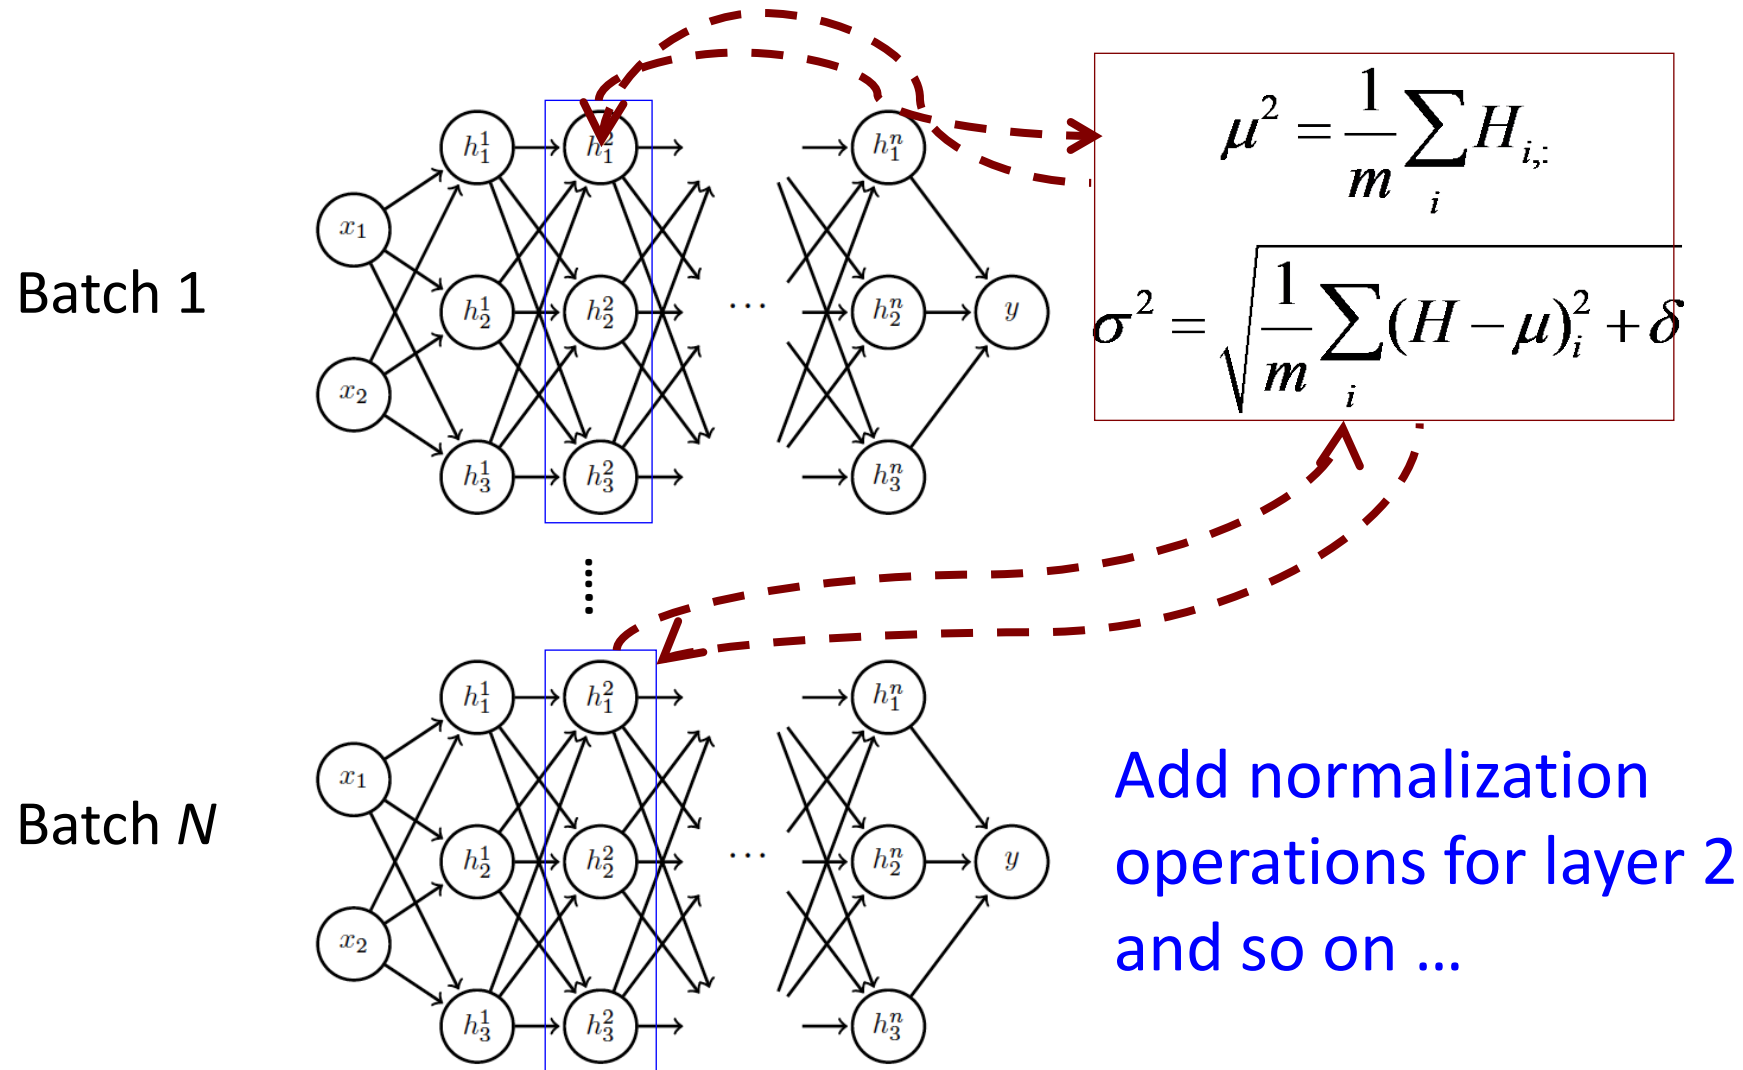
Learnable parameters

Allows network to **control range of normalization**

Batch Normalization



Batch Normalization



Regularization with unlimited computation

Best way to regularize a fixed size model is:

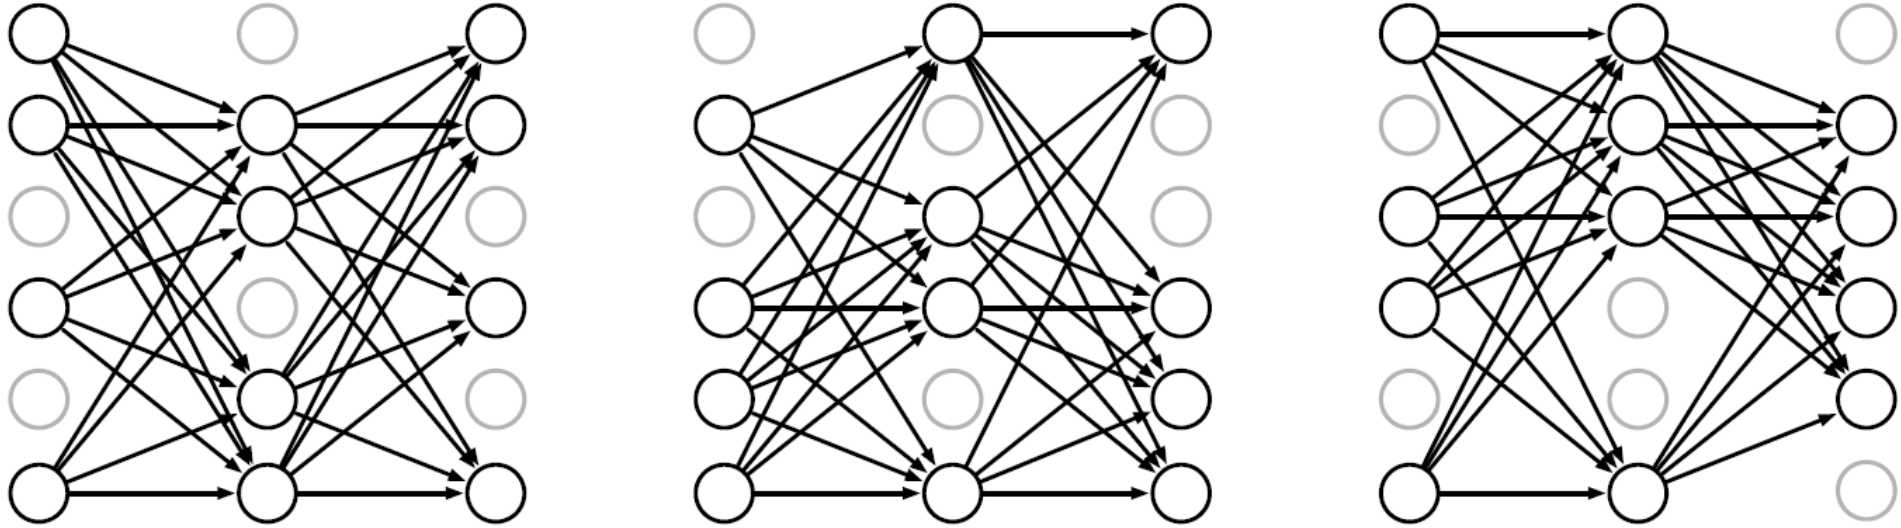
- Average the predictions of all possible settings of the parameters
- Weighting each setting with the posterior probability given the training data (Bayesian approach)

Dropout

Dropout does this using considerably less computation!

By approximating an equally weighted geometric mean of the predictions of an exponential number of learned models that share parameters

Dropout



- For each batch, different random set of nodes is removed
- Their values are set to 0 and their weights are not updated
- 10%, 20% or even 50% of all the nodes

Dropout is a bagging method

Bagging is a method of averaging over several models to improve generalization

Impractical to train many neural networks since it is expensive in time and memory.

Dropout makes it practical to apply bagging to very many large neural networks

It is a method of bagging applied to neural networks

Dropout is an inexpensive but powerful method of regularizing a broad family of models

Shortcuts and Highways

Deep learning: many layers of processing. Error propagation has to travel farther

All parameters in processing change have to be adjusted

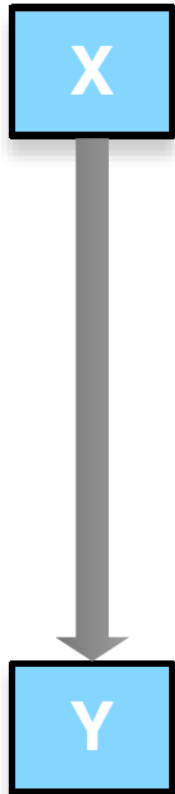
Instead of always passing through all layers, add connections from first to last

Jargon alert:

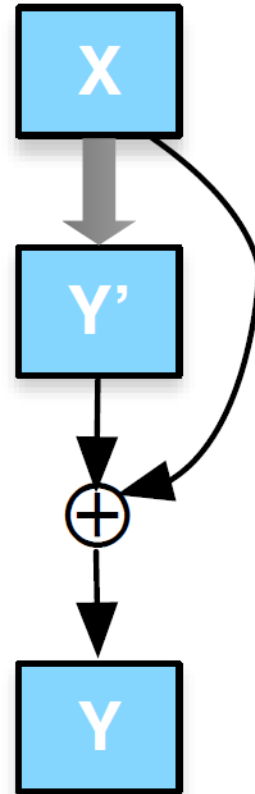
- Shortcuts
- Residual connections
- Skip connections

Shortcuts and Highways

Basic Layer



Skip Connection



Highway Network

